

{mthree}

Prometheus & Grafana

Site Reliability Engineering



Overview

In this module, we will look at a monitoring system called Prometheus, the dashboard called Grafana, and how they interact.

Learning Objectives

- Explain what Prometheus is and what it does
- Examine how Prometheus gets data
- Configure Prometheus
- Query Prometheus with PromQL
- Use Grafana to create visualisations

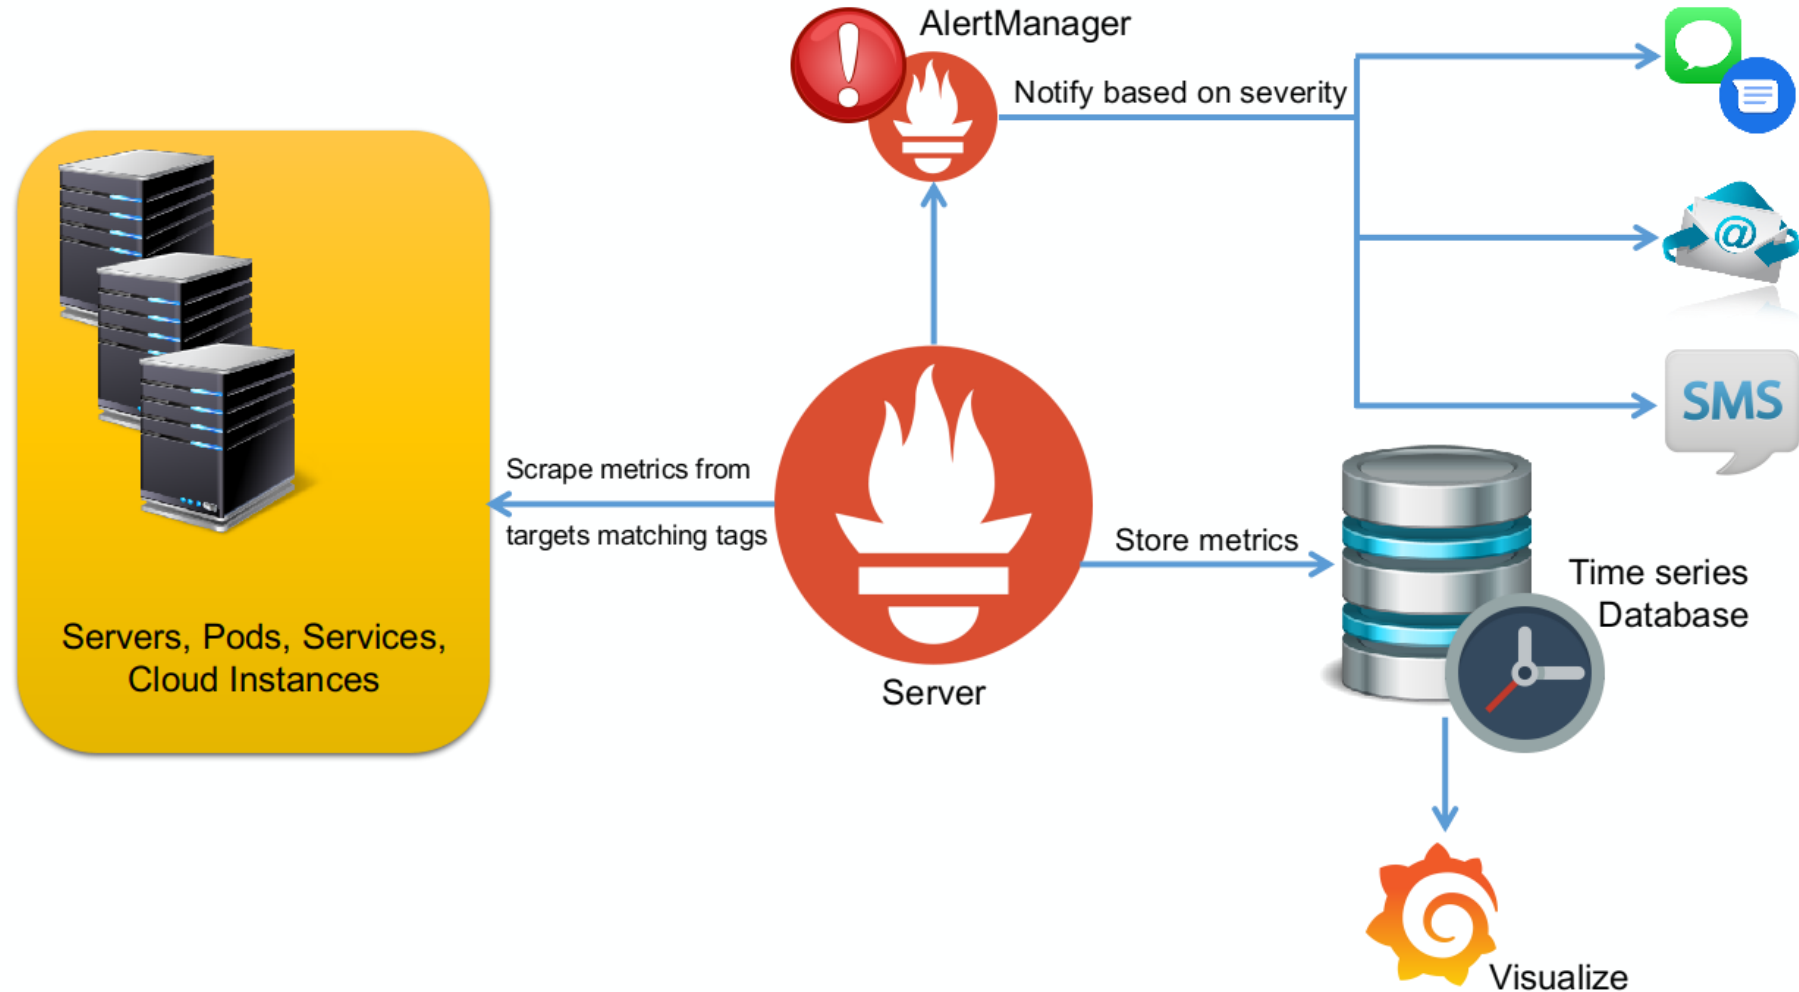


What is Prometheus?

- ↘ Monitoring tool
- ↘ From SoundCloud
 - >>> Created to monitor Highly dynamic environments
 - ~ Very useful for monitoring Kubernetes, Docker swarm
 - >>> Can be used in traditional non-container environments
- ↘ Has become the go to in Container and Microservice world
 - >>> Modern DevOps complex needs more and more automation
 - >>> Challenging and complex infrastructures need to be managed



How Prometheus Works



How Does Prometheus Work?

- ↘ Prometheus server
 - >>> Time series database
 - ~ Stores metrics data current CPU usage
 - >>> Data Retrieval worker
 - ~ Responsible for pulling metrics
 - >>> Web server
 - ~ Accepts queries and display data in UI or dashboard or visualization such as Grafana
- ↘ Monitors targets
 - >>> Linux server, Windows server, web servers, application servers, databases
- ↘ Measures units such as
 - >>> CPU status, request counts and durations, exception counts, memory/disk usage – Metrics



Monitoring Applications

- ↘ Collects relevant metrics such as:
 - >>> How many requests
 - >>> How many exceptions
 - >>> How many server resources
- ↘ Client libraries expose endpoints via HTTP URL called /metrics
 - >>> Requires using the library within your code
 - >>> <https://prometheus.io/docs/instrumenting/clientlibs/>
- ↘ Metrics can be
 - >>> Counter
 - >>> Guage
 - >>> Histogram



Exposing Metrics Steps

↘ Exporters

- >>> These are agents that extract data
- >>> Run on the system to monitor
- >>> Prometheus configured to point at the exporter agent port

↘ Many 3rd party exporters

- >>> <https://prometheus.io/docs/instrumenting/exporters/>
- >>> Pre-built - <https://prometheus.io/download/>

↘ Download and run exporter

- >>> For Linux, server download a node exporter, untar and execute
- >>> Converts metrics of the server
- >>> Exposes endpoints
- >>> Configure Prometheus to scrape this endpoint

↘ Windows is also supported

↘ There may already be a Docker container

- >>> Search the web for the app name and exporter
- >>> Or check through Prometheus exporters and see if they mention Docker containers

What do metrics look like?

- <https://jenkins.computerlab.online/prometheus/>



Configuring Metrics

- ↘ Exporters expose attributes
- ↘ Customize the exporters to define attributes to gather
 - >>> Generally passed as CLI options
 - >>> May have a configuration file
- ↘ Example MySQL metrics
 - >>> https://github.com/prometheus/mysql_exporter

```
./mysql_exporter -collect.auto_increment.columns \  
-collect.info_schema.innodb_metrics \  
-collect.perf_schema.tablelocks
```



Configuring Metrics for Java Spring Boot: Maven

↘ Maven pom.xml requires the following dependencies

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
<dependency>
    <groupId>io.micrometer</groupId>
    <artifactId>micrometer-registry-prometheus</artifactId>
    <scope>runtime</scope>
</dependency>
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-aop</artifactId>
</dependency>
```

Configuring Metrics for Java Spring Boot: `application.properties`

↘ `application.properties` requires

```
management.endpoints.web.base-path=/
management.endpoints.web.path-mapping.prometheus=/metrics
management.endpoints.web.path-mapping.health=status
management.endpoints.web.exposure.include=health,info,prometheus
management.metrics.distribution.percentiles-
histogram.http.server.requests=true
```



Configuration

➤ Prometheus.yaml – specify targets

- >>> How often
- >>> Rules for aggregation
- >>> Creating alerts
- >>> What resources to monitor
- >>> Define other endpoints

➤ Example shows static configuration

```
global:
  scrape_interval:      15s  evaluation_interval: 15s
.....
scrape_configs:
  - job_name: 'prometheus-metrics'
    static_configs:
      - targets: ['localhost:9090']
  - job_name: 'database'
    static_configs:
      - targets: ['mydb.courselab.online:9090']
```

Our Lab Configuration (1)

- Automated/Dynamic
- Prometheus scrapes Kubernetes pods and services
 - >>> Done through Custom Resource Definitions (CRD)
- Uses Prometheus Operator
 - >>> <https://prometheus-operator.dev/>
- Pods and Services expose scraper metrics ports
- New containers automatically detected
- PodMonitor or ServiceMonitor
 - >>> Depending on whether you are scaling

```
apiVersion: monitoring.coreos.com/v1
kind: PodMonitor
metadata:
  name: helm-controller
  namespace: flux-system
spec:
  selector:
    matchLabels:
      app: helm-controller
  podMetricsEndpoints:
    - port: http-prom
```



Our Lab Configuration (2)

- ↘ Managed via Kubernetes manifests
 - >>> Within your sre-course-infra environment/namespace directory
- ↘ Requires
 - >>> podMonitor.yaml or serviceMonitor.yaml or probe.yaml (if exists)
 - ~ To tell Kubernetes about your monitoring ports to scrape
 - >>> Containers
 - ~ Either a separate monitoring container, e.g. mysqld-exporter
 - ~ Or a port in your application container
 - That is not the application
 - Exposes Prometheus metrics



Example podMonitor

- Monitoring MySQL server remotely
- Reduce load on the Database server
- Dedicated pod to monitor



The Deployment Of Exporter

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/component: mysqld-exporter
    app.kubernetes.io/name: orderbook
  name: mysqld-exporter
  namespace: sre-example-dev
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/component: mysqld-exporter
      app.kubernetes.io/name: orderbook
  template:
    metadata:
      labels:
        app.kubernetes.io/component: mysqld-exporter
        app.kubernetes.io/name: orderbook
```

```
spec:
  containers:
    - env:
        - name: DATA_SOURCE_NAME
          value:
            "root:secret123@(orderbookdb:3306)/"
      image: prom/mysqld-exporter
      name: mysqld-exporter
      imagePullPolicy: Always
      ports:
        - name: http-prom
          containerPort: 9104
      restartPolicy: Always
```



The Prometheus Capture

```
apiVersion: monitoring.coreos.com/v1
kind: PodMonitor
metadata:
  name: mysql-exporter
  namespace: sre-example-dev
spec:
  selector:
    matchLabels:
      app.kubernetes.io/component: mysql-exporter
      app.kubernetes.io/name: orderbook
  podMetricsEndpoints:
    - port: http-prom
```

The port name for your
pod port exposing
/metrics



Monitoring the application using blackbox exporter

```
apiVersion: monitoring.coreos.com/v1
```

```
kind: Probe
```

```
metadata:
```

```
  name: orderbook
```

```
  namespace: sre-example-dev
```

```
spec:
```

```
  jobName: blackbox
```

```
  interval: 60s
```

```
  scrapeTimeout: 30s
```

```
  module: http_2xx
```

```
  prober:
```

```
    url: prometheus-blackbox-exporter.prometheus:9115
```

```
    scheme: http
```

```
    path: /probe
```

```
  targets:
```

```
    staticConfig:
```

```
      static:
```

```
        - http://orderbookdev.computerlab.online
```

Scraper URL



URL to monitor



Orderbook has /metrics

<https://cXXteam??dev.computerlab.online/metrics>

```
# HELP hikaricp_connections_usage_seconds Connection usage time
# TYPE hikaricp_connections_usage_seconds summary
hikaricp_connections_usage_seconds_count{pool="HikariPool-1",} 771241.0
hikaricp_connections_usage_seconds_sum{pool="HikariPool-1",} 9204.837
# HELP hikaricp_connections_usage_seconds_max Connection usage time
# TYPE hikaricp_connections_usage_seconds_max gauge
hikaricp_connections_usage_seconds_max{pool="HikariPool-1",} 0.0
# HELP hikaricp_connections_pending Pending threads
# TYPE hikaricp_connections_pending gauge
hikaricp_connections_pending{pool="HikariPool-1",} 0.0
# HELP jvm_gc_live_data_size bytes Size of long-lived heap memory pool after reclamation
# TYPE jvm_gc_live_data_size bytes gauge
jvm_gc_live_data_size_bytes 4.318788E7
# HELP system_cpu_count The number of processors available to the Java virtual machine
# TYPE system_cpu_count gauge
system_cpu_count 1.0
# HELP tomcat_sessions_alive_max_seconds
# TYPE tomcat_sessions_alive_max_seconds gauge
tomcat_sessions_alive_max_seconds 0.0
# HELP jvm_buffer_memory_used_bytes An estimate of the memory that the Java virtual machine is using for this buffer pool
# TYPE jvm_buffer_memory_used_bytes gauge
jvm_buffer_memory_used_bytes{id="mapped",} 0.0
jvm_buffer_memory_used_bytes{id="direct",} 81920.0
# HELP jvm_memory_committed_bytes The amount of memory in bytes that is committed for the Java virtual machine to use
# TYPE jvm_memory_committed_bytes gauge
jvm_memory_committed_bytes{area="heap",id="Tenured Gen",} 7.1983104E7
jvm_memory_committed_bytes{area="nonheap",id="CodeHeap 'profiled nmethods'",} 2.5952256E7
jvm_memory_committed_bytes{area="heap",id="Eden Space",} 2.883584E7
jvm_memory_committed_bytes{area="nonheap",id="Metaspace",} 7.23968E7
jvm_memory_committed_bytes{area="nonheap",id="CodeHeap 'non-nmethods'",} 2555904.0
jvm_memory_committed_bytes{area="heap",id="Survivor Space",} 3604480.0
jvm_memory_committed_bytes{area="nonheap",id="Compressed Class Space",} 9306112.0
jvm_memory_committed_bytes{area="nonheap",id="CodeHeap 'non-profiled nmethods'",} 1.540096E7
# HELP jdbc_connections_idle Number of established but idle connections.
# TYPE jdbc_connections_idle gauge
jdbc_connections_idle{name="dataSource",} 10.0
# HELP hikaricp_connections_creation_seconds_max Connection creation time
# TYPE hikaricp_connections_creation_seconds_max gauge
hikaricp_connections_creation_seconds_max{pool="HikariPool-1",} 0.0
```

Grafana

- ↘ Grafana allows you to:
 - >>> Query
 - >>> Visualize
 - >>> Alert on
 - >>> Understand metrics
- ↘ Central visualization of all your Prometheus monitoring
 - >>> Graphs, meters and more
- ↘ Dashboards to provide ease of understanding and speed
 - >>> Useful and colourful representation
 - >>> Thresholds and colours to determine issues or perfect conditions



Activity: Viewing Prometheus

Let's get a feel for Prometheus and PromQL

- ↘ Go to <https://prometheus.computerlab.online>
 - >>> Our Prometheus is authenticated through OAuth and your GitHub login
- ↘ Find the following:
 - >>> Jenkins disk usage
 - ~ Start by typing in disk into the search box
 - ~ Select the default_jenkins_disk_usage_bytes
 - ~ Click the Execute button
 - ~ You will now see various lines relating to bytes for various directories
 - >>> Typing any letter will bring up potential metrics that have been scraped
 - ~ e.g., probe_http
 - >>> Find the following:
 - ~ Kubernetes build info?
 - ~ What plugins do we have with Grafana?
 - ~ What namespaces exist in Kubernetes?



Activity: Viewing Prometheus

- View the pod container info for the orderbook container

```
kube_pod_container_info{container=~'orderbook.*'}
```

- Now look for your container orderbook in your namespace

```
kube_pod_container_info{container=~'orderbook.*',namespace='sre-example-dev'}
```


Exact match

Regex match


```
kube_pod_container_info{container="orderbookdb", container_id="docker://31c0d730d0fc3b37a21b9e7b50b4211b05d8ccd39ff48c7dc712403f515cf5fa", endpoint="http", image="108174090253.dkr.ecr.us-east-1.amazonaws.com/sre-course:orderbookdb-dev-31", image_id="docker-pullable://108174090253.dkr.ecr.us-east-1.amazonaws.com/sre-course@sha256:109ca209b63198abff412a9065958e32dfd714806a752ec4b8cc5f96be097e1a", instance="172.16.21.148:8080", job="kube-state-metrics", namespace="sre-example-dev", pod="orderbookdb-5d9577b9b4-xjr8k", service="prometheus-kube-state-metrics"}
```

Activity: PromQL in Grafana

- Head over to <https://grafana.computerlab.online>
- Click the Explore icon 
- Select Prometheus in the pull down, top left of screen



- In the metrics text field type your PromQL





- Enter kube_namespace_created; note IntelliSense as you type



Activity: PromQL in Grafana

- ↘ No results?
- ↘ Check the date range top right, left of Run Query
 - >>> Change to Last 7 days
- ↘ Have Grafana update
 - >>> Click the down arrow on the Run Query button
 - >>> Select a refresh interval
- ↘ Now query the Jenkins disk usage
- ↘ Change Range under Metrics to Instant
 - >>> What do you notice?
 - >>> Change back to Range
- ↘ Change the Step to 2d for every 2 days




Activity: Grafana Ui

- ↘ Take a look at the following;
 - >>> Click the Inspector - what does it tell you?
 - >>> Click the Query History - re-run the kube_namespace_created from here




Activity: Dashboards

- ↘ Dashboards provide your teams useful view of the site
- ↘ Creating through Grafana:
 - >>> Select the Dashboards icon 
 - >>> Select Manage from the pop-up menu
 - >>> From here you can create a new dashboard
 - ~ Create one with your full name for practice
 - ~ We can delete it later




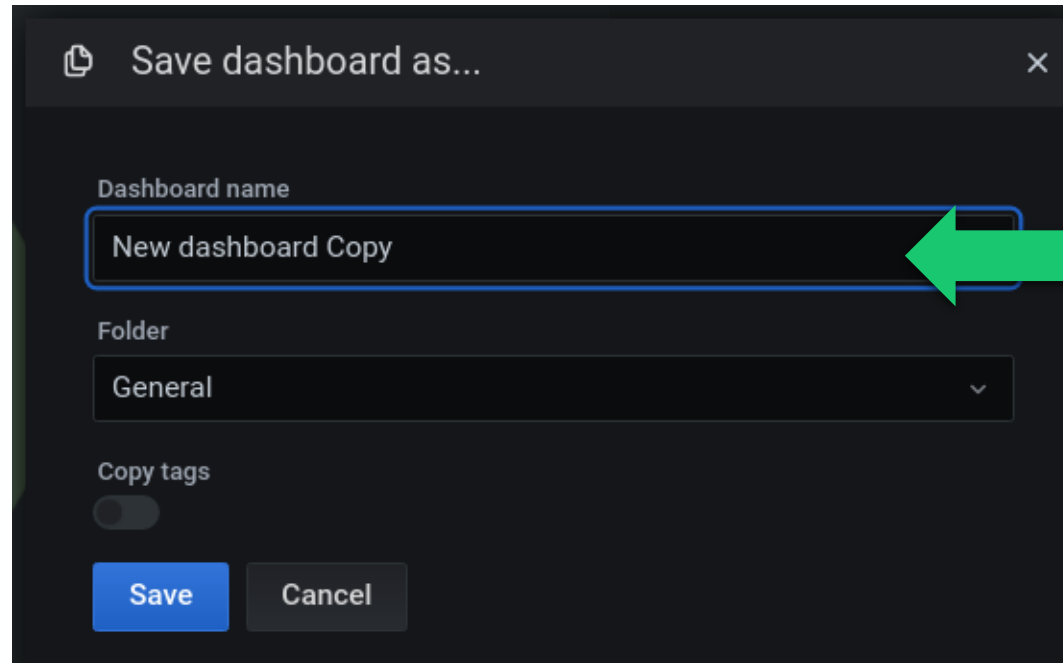
Activity: Dashboards

- ↘ Click the New Dashboard button
- ↘ You will be presented with a blank dashboard
- ↘ Click the Add new panel button to start adding charts
 - >>> Or use the icon 
- ↘ To the right of the screen are the settings for the panel
 - >>> Title your panel Jenkins Disk Space
 - >>> Expand the Visualization section
 - >>> Select a Pie Chart
 - >>> Under options
 - ~ Remove the Legend
- ↘ Add your metric to the Prometheus query in the lower main screen
- ↘ Select Instant
- ↘ Select Last hour as the time range
- ↘ Click Apply button to see your panel on the dashboard




Activity: Save the Dashboard

- Click the floppy disk icon to save 
- Fill in the pop up box by naming your dashboard
- If folders exist, you may wish to save it in a project folder



Activity: Folder

- ↘ If you wish to separate your work
- ↘ Click the + icon
- ↘ Select Folder
- ↘ Type in the name of your folder
- ↘ You can move existing Dashboards to the folder
 - >>> Go to Dashboards Home
 - >>> Select your Dashboard
 - >>> Click the settings icon 
 - >>> Select your folder from the pull down menu
 - >>> Save dashboard



Activity: Viewing Logs

- ↘ Our Grafana also trawls logs using Loki
 - >>> <https://grafana.com/oss/loki/>
 - >>> Alternatives include ELK, Splunk, Datadog
- ↘ Go to Explore
- ↘ Set the Explore data from Prometheus to Loki
- ↘ Where Metrics Browser used to be you'll now see Log browser
 - >>> You can click through to find a log
 - >>> You can start to type a log query



Activity: Show deployment logs

- ↘ Click Log labels
- ↘ Select namespace
- ↘ Select your Dev namespace
 - >>> If your namespace does not show in the list increase the time range
 - >>> Grafana only shows data in valid time range
- ↘ You will now see log lines from the files of all your running containers



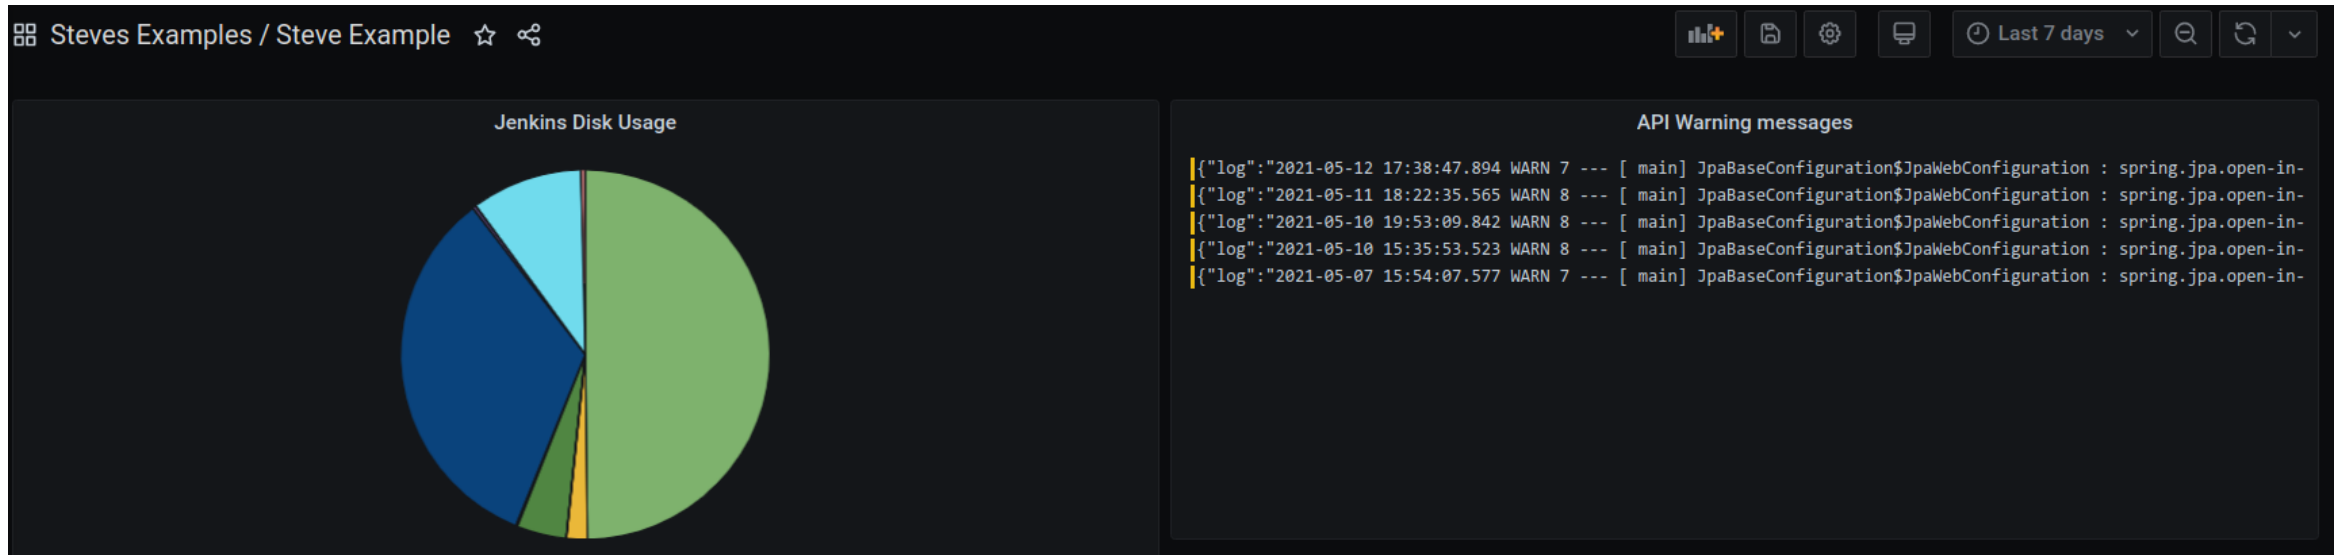
The Loki Query Syntax

- ↘ Similar to Prometheus
- ↘ Key = value statement contained in { }
- ↘ Exact match of value use =
- ↘ RegEx match of value use =~
 - >>> Note if you wish to "contain" then use .* either side of your text
 - ~ e.g. {namespace=~".*example.*"}
- ↘ Multiple key searches are comma separated
 - >>> {namespace="sre-example-dev", app="orderbookapi"}
- ↘ Values must be in double quotes
- ↘ Looking for the word error in the log line
 - >>> {namespace="sre-example-dev", app="orderbookapi"} |= "error"



Activity: Create A Log Entry

- Add to your dashboard a log entry for warnings in your orderbookapi
- Note you can drag widget positions around using the title bar
- The title bar also has a pull down menu next to the panel name



Activity: Coding the Dashboard 1

↘ Click the New Dashboard button

↘ Click the Add new panel button

>>> Set the panel properties:

~ Panel title: Application Memory Usage

~ Visualization: Gauge

~ Fields:

- Min: 0

- Max: 100

>>> Expand Thresholds and click Add threshold

~ 80 red

~ 70 yellow

↘ Set the query of Prometheus

>>> `sum(container_memory_usage_bytes{namespace="<yournamespace>", container="orderbookapi"}) / sum(container_memory_usage_bytes{namespace="<yournamespace>"}) * 100`

>>> Click the Apply button



Activity: Coding the Dashboard 2

↘ Click the Add new panel button

>>> Set the panel properties:

~ Panel title: Application CPU Usage

~ Visualization: Gauge

~ Fields:

- Min: 0

- Max: 100

>>> Expand Thresholds and click Add threshold

~ 80 red

~ 70 yellow


↘ Set the query of Prometheus

>>> `sum(container_cpu_usage_seconds_total{namespace="<yournamespace>", container="orderbookapi"}) / sum(container_cpu_usage_seconds_total{namespace="<yournamespace>"}) * 100`

>>> Click the Apply button



Activity: Export the Dashboard

- ↘ Click the  button
- ↘ Click JSON Model
- ↘ Click into the edit box and copy all of the JSON
- ↘ Save this to a file or other safe place for later use



Activity: Add the Dashboard Definition

↘ Switch to your local sre-course-infra repository

```
>>> git checkout main
>>> git pull
>>> git checkout <your branch>
>>> git merge main
```

↘ Copy an existing dashboard definition file to your project folder

```
>>> cd <project folder>
    ~ flux/apps/eks-sre-course/cXXX-team??-dev
>>> Edit the file status.yaml
```

↘ Edit stats.yaml and change:

```
>>> name: to cXXXteam??dev
>>> data: to cXXXteam??dev.json: |
    {
      <paste in the saved json definition>
    }
>>> At the bottom of the JSON code change
    ~ Change the title value to "cXXXteam??dev Dashboard"
    ~ Change uid value to "cXXXteam??dev"
```

↘ Save the new definition



Activity: Deploy the Dashboard Definition

- ↘ Push the changes to the shared repo
 - >>> git add --all (two dashes)
 - >>> git commit -m "Adding orderbook stats dashboard"
 - >>> git commit
 - >>> git push
- ↘ Log in to the shared repository
 - >>> Change to your branch
 - >>> Create a pull request
 - >>> Have the pull request approved
- ↘ The dashboard definition will be deployed



References & Further learning

- ↘ How Prometheus monitoring works
>>> <https://www.youtube.com/watch?v=h4SI21AKiDg>
- ↘ How To Setup A Grafana Dashboard Step By Step
>>> https://www.youtube.com/watch?v=4qpl4T6_bUw
- ↘ How to create Grafana Dashboards: The Easy way
>>> https://www.youtube.com/watch?v=Mqt_bBsejKQ
- ↘ Grafanalib: Dashboards as Code
>>> <https://www.youtube.com/watch?v=OOyEGG98B7w>
>>> <https://www.weave.works/blog/grafana-dashboards-as-code/>





Summary Q&A



Useful Prometheus Selectors

↘ probe_http_*

- >>> HTTP request information for the application
- >>> Provided you've added a probe to your namespace

↘ default_jenkins_*

- >>> default_jenkins_builds_duration_milliseconds_summary_sum
- >>> default_jenkins_builds_health_score
- >>> default_jenkins_builds_failed_build_count
- >>> default_jenkins_builds_success_build_count
- >>> default_jenkins_builds_last_build_result

↘ container_*

↘ kube_pod_*

↘ kube_deployment_*

↘ kube_endpoing_*

↘ kube_namespace_*

↘ mysql_*

↘ nginx_ingress_controller_*

↘ probe_http_*

↘ probe_success



Useful Loki LogQL

- ✚ `{app="kustomize-controller"}`
 - >>> To see if your deployment is there, failing, or changed
 - >>> add `|= "namespace"`
 - ~ To view only your namespace
- ✚ `{app="ingress-nginx"} | json | host="orderbookdev.computerlab.online"`
 - >>> To view request/response details to a particular applications
 - ~ orderbookdev.computerlab.online is the application
- ✚ `{app="ingress-nginx"} | json | host=~"orderbookdev.*" | request_time > 0.005`
 - >>> Checking if request time is > 5ms
- ✚ `sum(count_over_time(({container="orderbookac"} |= "Steve" |= "buy")[1h]))`
 - >>> The number of buy's by Steve over the last 1 hour period
- ✚ `sum(rate(({container="orderbookac"} |= "Steve" |= "sell")[1m]))`
 - >>> The per second rate of all sell's within the last minute
- ✚ `{namespace="orderbook-dev"}`
 - >>> Get logs from containers running in your namespace

