



Kubernetes

Site Reliability Engineering



Overview

Learning Objectives

In this module, you will be introduced to the deployment system

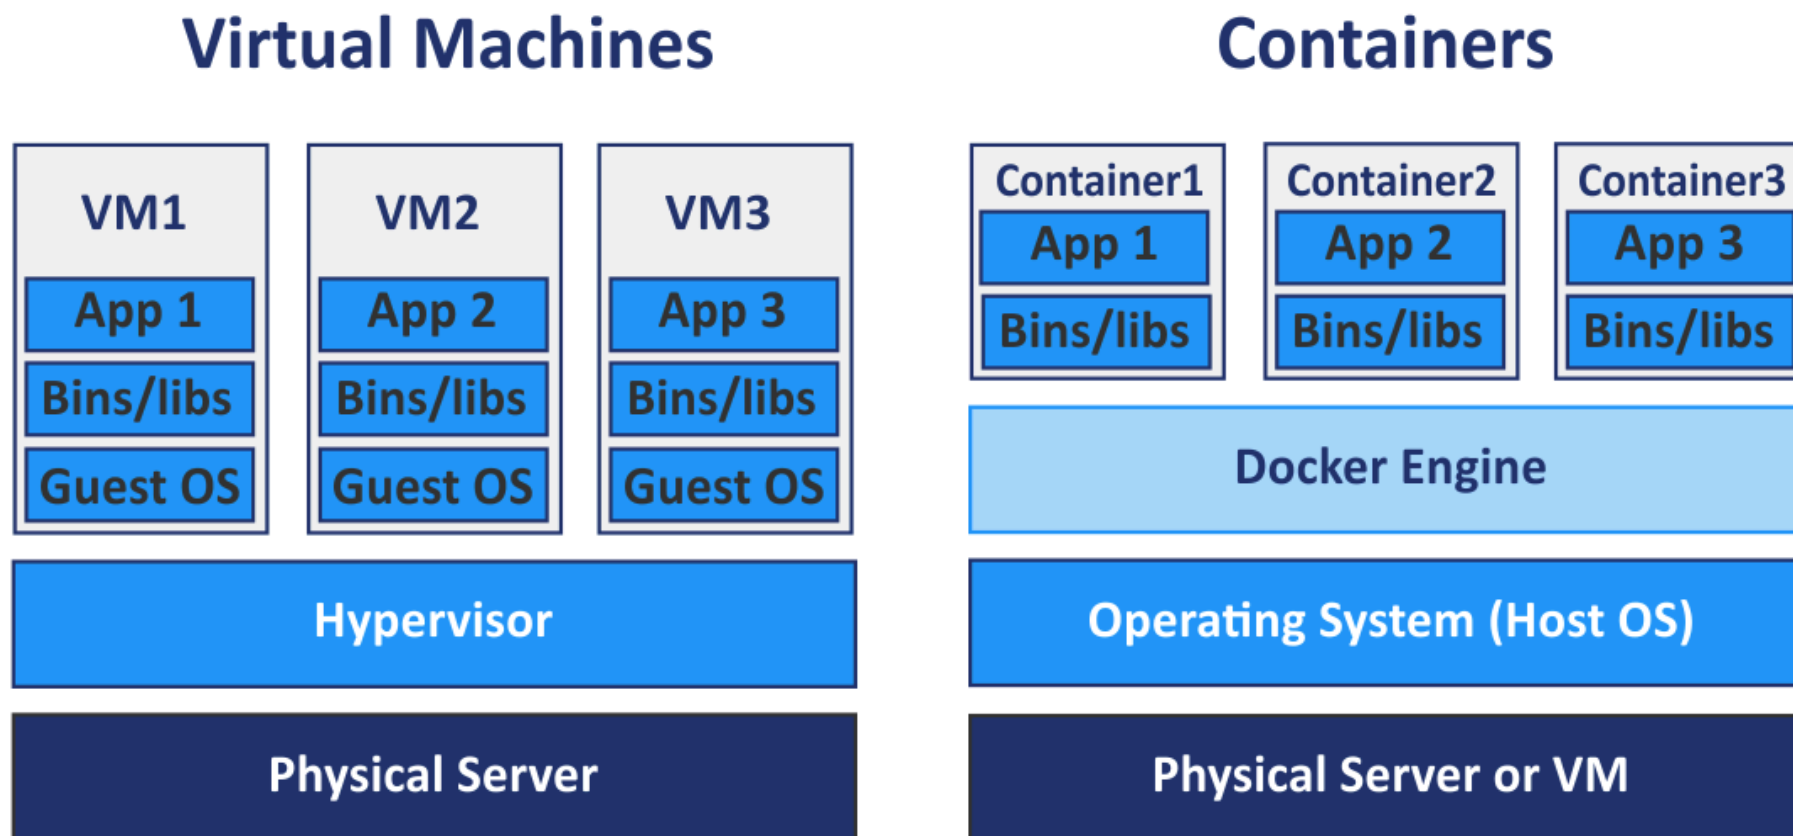
By the end of this module, you will be able to:

- ↘ Explain what a container service is
- ↘ Review and modify infracode
- ↘ Debug deployment issues
- ↘ Explain the difference of create vs. update

Kubernetes

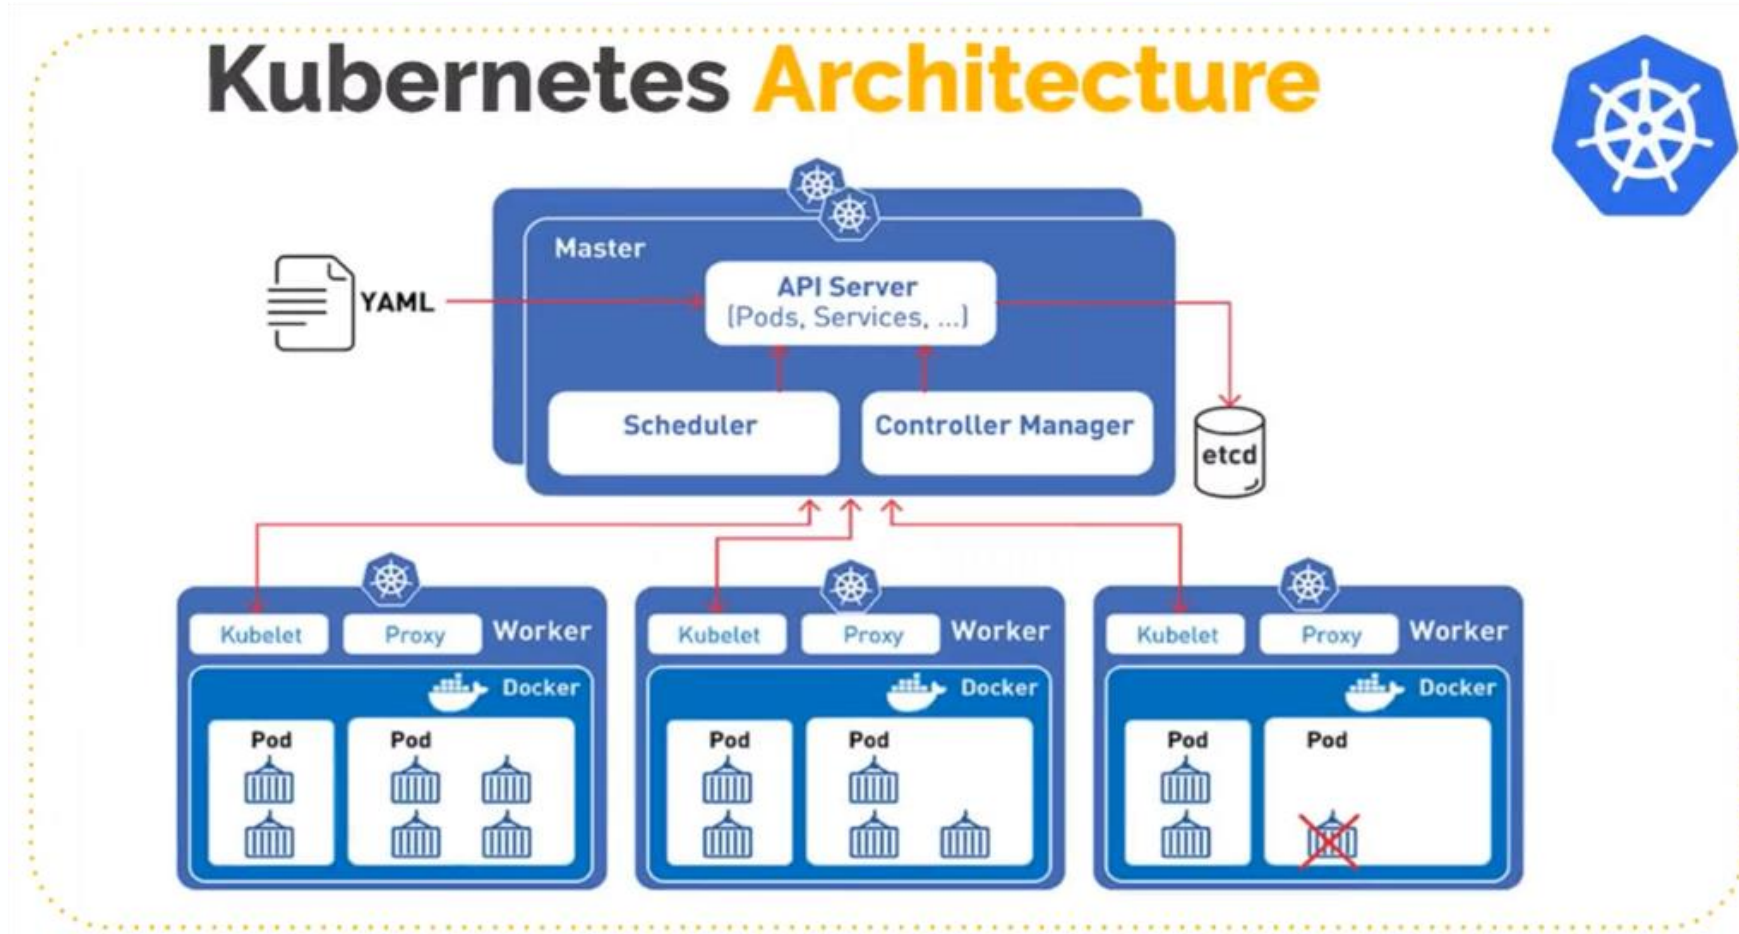
- ↘ A container orchestration tool to manage enterprise scale microservices
- ↘ Manages
 - >>> Container infrastructure (networking, load balancing)
 - >>> Deployments
 - >>> Persistent Storage
 - >>> Secrets
 - >>> Variables

VMs vs Container Deployments



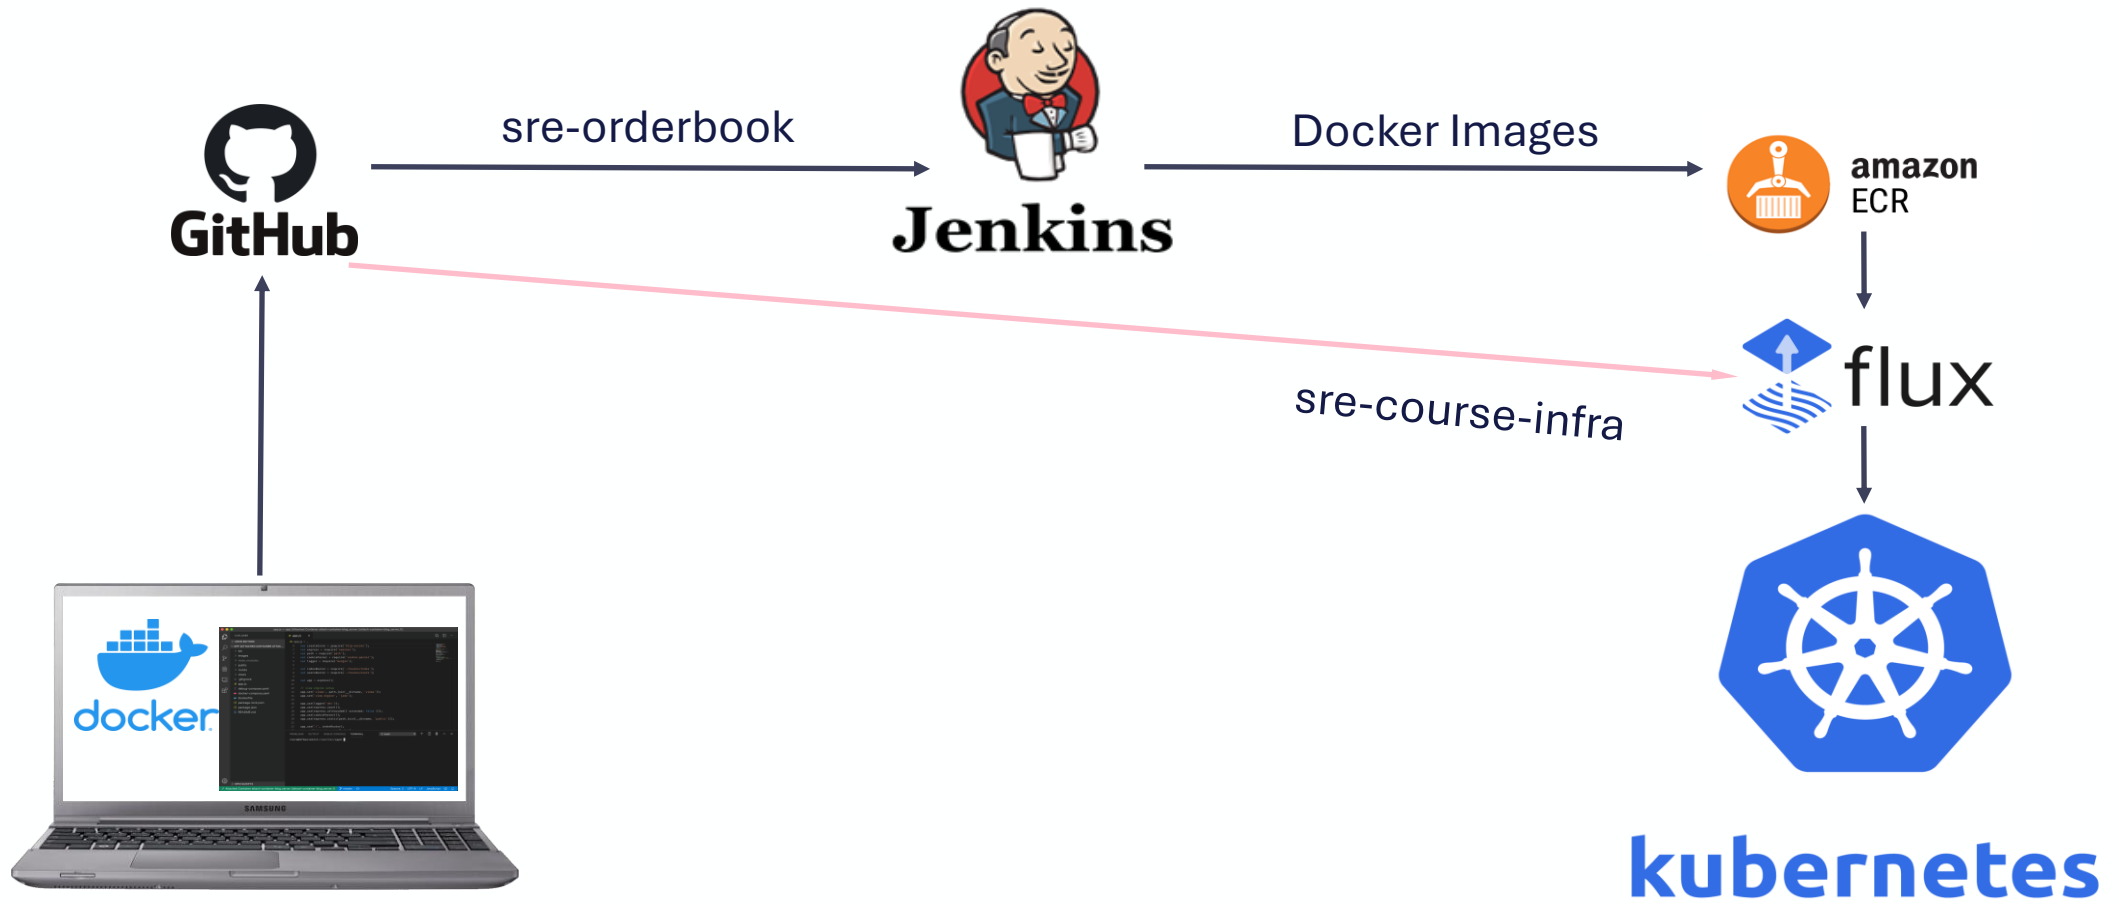
Source: [Docker containers are not lightweight virtual machines](#)

Containers in Kubernetes




Source: [Getting Started With Kubernetes - Part Two](#)

Containers to Kubernetes



Our Lab Environment



Services of Help

↘ Jenkins

>>> <https://jenkins.computerlab.online>

↘ Grafana Monitoring Dashboard

>>> <https://grafana.computerlab.online>

↘ Prometheus metric gathering

>>> <https://prometheus.computerlab.online>

↘ Alertmanager

>>> <https://alertmanager.computerlab.online>

↘ Docker registry repository

>>> <http://ecrlist.computerlab.online/index.php>

↘ Kubernetes

>>> <https://k8sdashboard.computerlab.online>

The Lab Repository

- ↘ Kubernetes is managed through
 - >>> [GitHub.com/The-Software-Guild/sre-course-infra](https://github.com/The-Software-Guild/sre-course-infra)
- ↘ DevOps practices must be used with this repository
 - >>> Coding on feature branches
 - ~ Never code on main
 - >>> Pull requests required to merge your code to main branch

Activity: Kubernetes – Step-by-Step

IMPORTANT: Naming convention!!!!
Use lowercase always!!!!

↘ Create Dev and Production environments

- >>> Using the sre-course-infra Git repository and the make-sre-env script
 - ~ This is a one-off task for each environment
- >>> Make sure you use a branch to do your updates
 - ~ Branch name to follow the format of ?XXXteamNN
(? = r – reskill c – cohort)

↘ Deploy Dev

- >>> Use PR in Github web page to merge your changes to main
- >>> FluxCD will detect your changes to the sre-course-infra and create the environment, or will identify new Docker images in the ECR and deploy

↘ Promote to Production

- >>> Make a copy of the Hardcoded Promote Jenkins job
 - ~ Change orderbook in the environment section to your ?XXXteamNN Jenkins job name
(? = r – reskill c – cohort)
- >>> Modify to create your own Production image
- >>> Or view <https://github.com/The-Software-Guild/sre-course-examples>
 - ~ Branch d1m7t1
 - ~ Directory Day1/Module7/Task1/instructor-promote-prod
 - ~ Change instructor to your ?XXXteamNN Jenkins job name
(? = r – reskill c – cohort)

↘ Deploy to Production

↘ View the logs in Grafana to detect any errors

- >>> Use {job="flux-system/kustomize-controller"} using Loki to determine if there are any errors in your sre-course-infra code

Activity: Creating Environments – Step-by-Step

This is a one-off task!

↘ Create Dev

>>> Clone the following repository

~ <https://github.com/The-Software-Guild/sre-course-infra>

>>> Create a new branch called ?XXX-teamNN
(? = r – reskill c – cohort)

~ Where XXX is your course code (instructor can help if you do not know)

~ NN is the team number assigned by your instructor

>>> Run the script called make-sre-env

~ You will need to supply the following as command line arguments

- Course in the format of ?XXX
(? = r – reskill c – cohort)
- Team in the format of teamNN
- Environment to create dev
- Example - ./make-sre-env c149 team01 dev

>>> Add, commit and push your branch

>>> Create a Pull Request to merge to the main branch

>>> Merge if PR allowed, or fix conflicts then come back to merge

↘ Create Prod

>>> In the same git repository run make-sre-env again, but change dev to prod

>>> Perform the same actions as above for getting your deployment actioned

IMPORTANT: Naming convention!!!!
Use lowercase always!!!!

Activity: Changes to Environment – Step-by-Step

- ✚ Any changes should be done on your branch
- ✚ Follow the add, commit, push, and pull request method to release
- ✚ Always check that there are no deployment errors in Grafana/Loki
- ✚ Where did make-sre-env put my manifests?
 - >>> **flux/apps/eks-sre-course/?XXX-teamNN-[dev|prod]**
 - ~ Use this location for any changes you wish to make to your application infrastructure
 - ~ All actions for Kubernetes to perform
 - e.g., deployment modification, DNS name for ingress, service names, etc

namespace.yaml

Created by the `make-sre-env` script

➤ This defines the unique place for all your application components in isolation

```
apiVersion: v1
kind: Namespace
metadata:
  name: sre-example-dev
```

Manifest format that the Kubernetes cluster API will need to interpret

The manifest type

Naming the namespace

Access from the WWW – ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: orderbook
  namespace: sre-example-dev
  annotations:
    kubernetes.io/ingress.class: "nginx"
    cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
  rules:
    - host: sreexampledev.computerlab.online
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: orderbookapi
                port:
                  number: 8080
```

Namespace to be applied to

External service that knows us;
Required for HTTPS connections

Domain name to access this
service

URL to get to the service

Service name and port
in Kubernetes

Manifests

↘ Define resources in Kubernetes

↘ Kinds of manifest

>>> Namespace

>>> Deployment

>>> Service

>>> Ingress

Activity: Checking the App

- ↘ The application:
<http://?XXXteamNNdev.computerlab.online>
- ↘ We can also check the status
 - >>> By clicking the **View App Status**
 - >>> By going to <http://?XXXteamNNdev.computerlab.online/status>
- ↘ Click **View Order History** link to view some preloaded transactions.

NOTE: Your URL can be found in the **ingress.yaml** file value of the **host** attribute


Home

[Submit a Buy Order](#)
[Submit a Sell Order](#)
[View Order History](#)
[View Portfolio](#)
[View App Status](#)

Activity: Checking the Logs

↘ Check that Prometheus and Grafana have picked up our environment

>>> Navigate to <https://grafana.computerlab.online>



↘ Here we want to Explore

>>> Click the pull-down to the right of the Explore button and select Loki if it is not already selected.

>>> Below this there are some query options

>>> Click the pull down called Log labels and select namespace

>>> Scroll down to find your namespace which should be named team??dev

↘ Click the Run Query button, top left

↘ You should see the logs relating to your namespace for the containers

Check for Deployment Errors

- ✚ Errors can break the build for everyone
 - >>> You need to check the logs
 - >>> Loki query `{job="flux-system/kustomize-controller"} |= "error"`
 - ~ This can be used to determine if there are errors in Flux deployments
- ✚ LogQL – used by Loki in Grafana
 - >>> [LogQL in Grafana Loki](#)

Activity: Kubernetes Failing Application

- ↘ Make a deployment fail if it hasn't
 - >>> We can do this by making a typo in the image name
 - ~ Edit the deployment-api.yaml file
 - ~ Locate image line
 - ~ Change *orderbookapi-dev* to *notebookweb-dev*
 - ~ Remove the entire comment following the image value
 - ~ Add, Commit, Push, PR merge
 - ~ Check k8sdashboard and grafana for errors, and alertmanager
- ↘ Review the failure in Grafana and Prometheus
- ↘ Return production to normal
 - >>> Remove the typo: Recreate from the api.yaml file
- ↘ Prevent failure and reliable deployments
 - >>> [Kubernetes deployment strategies](#)

Activity: View the Error

- Open <https://alertmanager.computerlab.online>
 - >>> This is our alerting system when problems occur
 - >>> If there is an issue with your pod launching, you will find it here
- In the filter, you can type things such as
 - >>> namespace="yourNamespace"
 - >>> container="yourContainername"

The interface shows a 'Filter' tab with a 'Group' dropdown. On the right, there are controls for 'Receiver: All', 'Silenced', and 'Inhibited'. A filter rule is entered: 'namespace="sre-example-dev"' followed by a red 'x' icon. To the right of the input is a blue '+' button and a 'Silence' link. Below the input field, a hint reads: 'Custom matcher, e.g. env="production"'.

The interface is identical to the one above, but with an additional filter rule. The first rule is 'namespace="sre-example-dev"' with a red 'x' icon. The second rule is 'container="orderbookapi"' with a green border and a red 'x' icon. The blue '+' button and 'Silence' link are still present. The hint text 'Custom matcher, e.g. env="production"' remains at the bottom.

Identify the Error

Alertmanager Alerts Silences Status Help

New Silence

Filter Group

Receiver: All ☐ Silenced ☐ Inhibited

+

Silence

Custom matcher, e.g. `env="production"`

+ Expand all groups

+

Not grouped1 alert

-

Not grouped1 alert

16:10:28, 2021-05-04 (UTC)

+ Info

Source

Silence

alername="KubeContainerWaiting"

+

container="orderbookapi"

+

namespace="sre-example-dev"

+

pod="orderbookapi-7dc698c895-zhlx9"

+

prometheus="prometheus/prometheus-kube-prometheus-prometheus"

+

severity="warning"

+

The problem

The container

Namespace

The pod

Alert: CrashLoop

"Application fails to run" generally causes this issue

Pod failing to start

17:39:58, 2021-05-04 (UTC) [+ Info](#) [Source](#) [Silence](#)

alertname="KubePodCrashLooping" +

container="orderbookapi" +

endpoint="http" +

instance="172.16.0.149:8080" +

namespace="sre-example-dev" +

pod="orderbookapi-94dc87868-6p77z" +

prometheus="prometheus/prometheus-kube-prometheus-prometheus" +

service="prometheus-kube-state-metrics" +

severity="warning" +

Deployment Manifests

Deployment kinds

- Define how your containers are to be deployed into the environment
- Name of the container/deployment
- Variables to be passed
- Ports to expose for the service to be used
- The image to be deployed
- How many to launch

References

- >>> [Understanding Kubernetes Objects](#)
- >>> [Kubernetes: Deployments](#)
- >>> [Kubernetes: Managing Resources](#)
- >>> [Kubernetes API](#)
- >>> [Kubernetes: Deployment](#)

Preventing Failure

- ↘ RollingUpdate
 - >>> Create a ReplicaSet for the deployment
 - >>> The ReplicaSet is tested (if applied) to ensure it works
 - ~ If the pod dies, the replicaset is destroyed
 - >>> The original ReplicaSet remains running until the new one works
 - ~ Allows existing connections to be completed (pool draining)
- ↘ Create
 - >>> Destroys the current ReplicaSet and then launches the new one
 - ~ All existing pods/containers are destroyed immediately
- ↘ Useful reading
 - >>> [Kubernetes deployment strategies](#)

Pros/Cons

Create

- ↘ Pros
 - >>> Application state entirely renewed
- ↘ Cons
 - >>> Downtime that depends on both shutdown and boot duration of the application

RollingUpdate

- ↘ Pros
 - >>> Version is slowly released across instances
 - >>> Convenient for stateful applications that can handle rebalancing of the data
- ↘ Cons
 - >>> Rollout/rollback can take time
 - >>> Supporting multiple APIs is hard
 - >>> No control over traffic

Blue/Green Deployment

- ↘ Create and RollingUpdate are pod-based
 - >>> Update allows for a slow roll out and protects current connections
- ↘ What if our old API will cause issues with the new roll out?
 - >>> For example, a missing field in the API when updating the database
- ↘ Blue/Green deployment
 - >>> Ensure the load balancer sends traffic only to the new version
 - >>> The service is updated and uses a version label in the spec.selector section
 - >>> The deployment will also need to have a version label for matching

Blue/Green Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-app
  labels:
    app: my-app
spec:
  type: NodePort
  ports:
    - name: http
      port: 8080
      targetPort: 8080

# Note here that we match both the app and the version.
# When switching traffic, we update the label "version" with the appropriate value, ie: v2.0.0
selector:
  app: my-app
  version: v1.0.0
```

Blue/Green Pros/Cons

Pros

- Instant rollout/rollback
- Avoid versioning issue, change the entire cluster state in one go

Cons

- Requires double the resources
- Proper test of entire platform should be done before releasing to production
- Handling stateful applications can be hard

Canary Deployment

- ↘ Running multiple versions at the same time
- ↘ Good for user testing
- ↘ Requires two deployments and more management
 - >>> Monitoring of traffic and pods required
 - >>> Need to know which users are working with which version
- ↘ Defined using a percentage split (e.g., 75% old – 25% new)
- ↘ Service uses **app** label selector, no version
- ↘ Deployments use same label for app
 - >>> Different version label
 - >>> Different number for **replicas**, e.g., 3 and 1

Canary Pros/Cons

Pros

- ↘ Version released for a subset of users
- ↘ Convenient for error rate and performance monitoring
- ↘ Fast rollback

Cons

- ↘ Slow rollout
- ↘ Fine tuned traffic distribution can be expensive
- ↘ More detail in monitoring
- ↘ Complex support
 - >>> Need to know which version users are using

Other Features – 1

The deployment is the key element

- ↘ Look at the features that the deployment offers
 - >>> Strategy
 - ~ Allows defining of the type of deployment, rolling, recreate
 - >>> Replicas
 - ~ The number of containers to start
 - >>> MinReadySeconds
 - ~ How long to wait to ensure container isn't crashing
- ↘ Pod template spec parameters
 - >>> Volumes
 - ~ Enable persistent storage
 - ~ [Configure a Pod to Use a Volume for Storage](#)
 - >>> RestartPolicy
 - ~ Whether to restart the container if it fails
 - >>> Liveness and Readiness probes
 - ~ Determine if the container is still running and ready to run using a status check
 - ~ [Configure Liveness, Readiness and Startup Probes](#)

Other Features - 2

↘ Secrets

- >>> Passwords
- >>> SSH Keys
- >>> SSL Certificates

↘ Add through

- >>> secretGenerator yaml files
 - ~ [Managing Secrets using Kustomize](#)

↘ A secret manifest

- >>> [Secrets](#)

```
secretGenerator:  
- name: db-user-pass  
  literals:  
  - username=admin  
  - password=1f2d1e2e67df
```

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: db-user-pass  
type: kubernetes.io/basic-auth  
stringData:  
  username: admin  
  password: 1f2d1e2e67df
```


Other Features – 3

- ↘ ConfigMaps
 - >>> Allow storing of variables and values for use in deployments
 - >>> Pods/Containers may require variables to be set
 - >>> [ConfigMaps](#)
- ↘ ServiceAccounts
 - >>> Used to allow pods to interact with Kubernetes services
 - >>> [Configure Service Accounts for Pods](#)



Summary Q & A