

Monitoring

Site Reliability Engineering



Overview

Learning Objectives

In this module, we will provide a high-level overview of why we need monitoring and the ways in which we use it.

By the end of this module, you will be able to:

- ↘ Explain what it means to monitor
- ↘ Describe why it is needed and its benefits
- ↘ Describe the different layers of monitoring and types
- ↘ Explain what distributed monitoring is
- ↘ Identify the different methods systems use

What is Monitoring?

- ↘ Think of something you do every day.
 - >>> Going for a run or to the gym
 - >>> Walking or driving to a new place
 - >>> Cooking a meal or cake
 - >>> Catching a bus or train

- ↘ Do you monitor anything while performing these tasks?



Yes, You Do

- ↘ Running or gym
 - >>> Monitor heart rate or blood pressure
 - >>> See if it is getting to the desired target
- ↘ Walking or driving to a new
 - >>> Check road signs or names to make sure you are going the right way
- ↘ Cooking a meal or cake
 - >>> Set timers
 - >>> Check weights and measures
 - >>> Check that it is cooked inside
- ↘ Catching a bus or train
 - >>> Check the platform and arrival time
 - >>> Do I have enough time to get a coffee?



What is Monitoring?

The collecting, processing, aggregating
and displaying of
real-time quantitative data about a
system.

Why Monitoring?

- ↘ We can make executive decisions from the data
 - >>> Increase resource
 - >>> Add more storage
 - >>> Design a better process
 - >>> Fix site and system flaws
- ↘ Report findings in a timely fashion
 - >>> Dashboards to display useful charts, data and diagrams
 - >>> Alerts to email, SMS or chat
 - >>> Escalate issues quickly
 - >>> Determine if an issue is really occurring
 - ~ e.g., Did my CPU just hit 100%? Is it still there?
- ↘ Perform trend analysis through historical data
 - >>> Helps with planning and delivery of reliable systems
 - >>> Foresee potential demand through external events such as political events, or disasters
 - >>> Link to machine learning and AI systems to predict future trends
- ↘ Compare behaviours between different versions of applications or system components



Activity: Setting a Baseline

- ↘ A key part of monitoring is setting a baseline.
- ↘ We need to understand our system before we can make assumptions
- ↘ Let's monitor our individual heart rate for the next minute.
 - >>> Record that number
- ↘ Now let's get our heart rate going a little
 - >>> Perform an above the head clap 20x as fast as you can
 - >>> Then measure your heart rate for the next minute


- ↘ In your teams, determine the average value and record the value
 - >>> Do this for both normal and after the above the head clap
 - >>> What is the variance between you all?
- ↘ Now across teams, what is the average value of BPM?
 - >>> Do this for both normal and after the above the head clap
 - >>> How big a difference is your team value to the average

Activity Outcome

- ↘ The individual value
 - >>> Getting an idea of each part of the entire system
 - ~ Disk, CPU, memory, network, traffic, processes, thread count and so on
- ↘ The average value in the team
 - >>> Equivalent to using the same rule for all systems
 - ~ e.g., Disk usage or CPU utilization
- ↘ Across the different teams
 - >>> Equivalent to using our rule for the whole organization
- ↘ Accelerated heart rate
 - >>> We need to understand where our peak load occurs


Baselines

- ↘ Give us the starting point of monitoring
 - >>> Normal activity
 - >>> Increased activity
 - >>> Low activity
- ↘ Obtaining values for production on new projects
 - >>> Testing as part of Dev
 - ~ Capacity and Performance
 - ~ Playing potential scenarios to see what results you get
 - ~ Recording the outcomes as thresholds
- ↘ Using values from previous projects
 - >>> Where we have "like" usage




Thresholds

- ↘ Baselines help us set thresholds
- ↘ Setting the limits and boundaries for
 - >>> Perfect operation
 - >>> Forewarning of potential problem
 - >>> Critical when the issue is impacting the client or system



Observability

- As systems evolve, so does our monitoring for business impact
- Observability in an industrial system example:
 - >>> Adding a flow gauge inside a water pipe
 - ~ Like the internals of an application
 - >>> Connecting it to an external display
 - ~ The telemetry – dashboard view
 - >>> Operator can observe internal property of the system through external device
 - ~ e.g. how fast water is flowing inside a pipe
 - >>> Threshold can be set to send alerts



Benefits of Monitoring

Business

- Reputation
- Keeping the user happy
- Audit compliance
- Discover user trends

Technical/Project

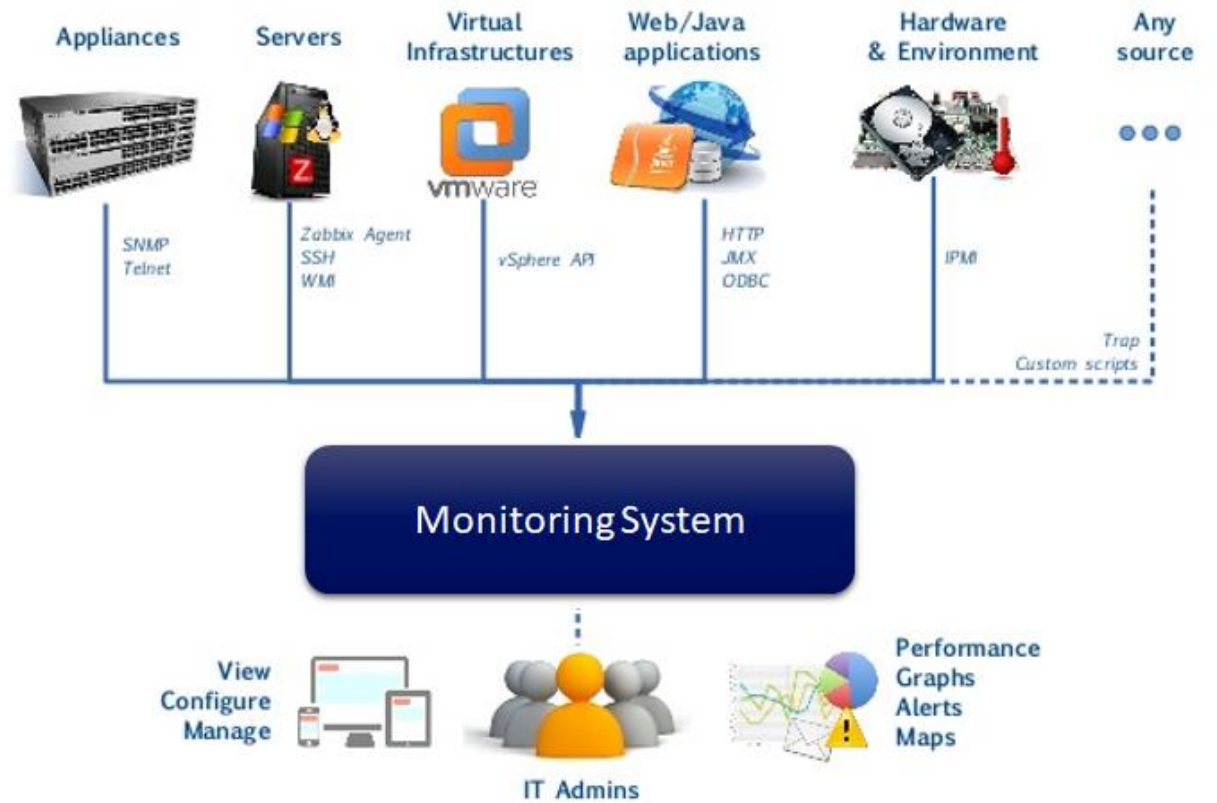
- Pre-empt failure
- Security threat detection
- Early problem detection
- Notification and visualization
- Performance analysis

Both

- Historical trend analysis and future prediction
- Meeting service agreements and objectives
- Planning and budgeting

Classic Monitoring

- ↘ System resources
 - >>> Hardware components
 - >>> Operating system
- ↘ Network components
 - >>> Switches
 - >>> Load balancers
 - >>> Firewalls
 - >>> Routers
- ↘ Applications
 - >>> Internals
- ↘ Metrics, rules and thresholds
- ↘ Reactive




2 • COLLECT AND VIEW DATA

Modern Systems

- ↘ Centralized dashboards, collection and alerting
- ↘ Gathers
 - >>> Metric data
 - >>> Log file information
 - >>> Either by requesting or receiving
 - >>> Application and infrastructure
- ↘ Queries, calculations & trend analysis
- ↘ Extensible and pluggable
 - >>> Additional information gathering and calculations
- ↘ Low impact agents
 - >>> Gathering data from resources being monitored



Push monitoring systems




ITRS/Zabbix

Configuration defined at the server
Data stored at the server and predefined



Configuration pushed to agent

Metrics sent to the server per period of time



Netprobe/Zabbix Agent
Special agent for monitoring



Pull monitoring systems



Prometheus

Data stored at the server
Raw data time series
Key/Value



Metrics requested via http(s) on a polling basis



Application defines metrics

Infrastructure vs Application Monitoring

Application

- >>> Underlying system components
 - ~ Operating System versions and packages
 - ~ Hardware – Disks, CPU, Memory
 - ~ Network components – firewalls, proxy, routers, etc
- >>> Interested in health, load, latency and errors in logs
- >>> Configuration changes or errors
- >>> Monitoring for potential failure based on state



Infrastructure

- >>> Dependencies to other services
 - ~ Connectivity, latency, availability
- >>> Internals
 - ~ Code performance and errors
 - ~ Function/method speed
- >>> Observability of how customer is seeing the application
 - ~ Technical and business metrics



Black-Box vs. White-Box Monitoring

White Box	Black Box
Modest Use	Critical Use
<ul style="list-style-type: none">• Application Monitoring• Internal metrics exposed<ul style="list-style-type: none">• Logs• HTTP endpoints• Java metrics<ul style="list-style-type: none">• JVM Profiling• JMX• Metrics such as:<ul style="list-style-type: none">• Number of http requests• SQL queries running	<ul style="list-style-type: none">• Server monitoring<ul style="list-style-type: none">• Standard metrics such as:<ul style="list-style-type: none">• Disk• CPU• Memory• Other metrics<ul style="list-style-type: none">• Network switch traffic/access• Load balancers• Hypervisor resource usage• Hard disk and other hardware errors• System-oriented problems<ul style="list-style-type: none">• Active, not predicted

Aggregate Metrics

- ↘ Aggregates must be
 - >>> Meaningful
 - >>> Relevant
- ↘ What time period is relevant?
 - >>> Is an hour too short – will we trigger an alert
 - >>> What's our objective to the client
- ↘ Am I measuring in the right place?
 - >>> User latency may be hard to identify
 - >>> Where can I measure latency?
 - ~ In bound proxy/load balancers
 - ~ Do I record at all network connected points and take the sum/average
 - >>> What if I cannot pin point a particular request
 - ~ Do I take the average of all traffic for the period?



Calculations

- ↘ Long term view of data
- ↘ Aggregate usually sufficient to facilitate growth planning
- ↘ Detailed individual metrics useful
 - >>> May be expensive
 - >>> Impractical to store and retrieve
- ↘ Use counters
- ↘ Use percentiles
 - >>> 50th, 95th and 99th percentiles lets you see 50%, 5% and 1% of requests are too slow
- ↘ If not available
 - >>> Mean value by summing the seconds spent in requests and dividing by the number of requests
 - >>> Logging every request and computing the percentile values by scanning or sampling the log entries

BUT do they tell the truth?

Monitoring Production

- ↘ Is it enough just to monitor production environment?
- ↘ What if the customer was expecting the update today?
- ↘ Is our pipeline functioning correctly?
 - >>> Are all the agents on line and available
- ↘ Is the QA environment set to the correct versions
 - >>> Do we have the correct infrastructure set up?
 - >>> Are the tests up to date?
- ↘ Are the correct versions in our software repository?
- ↘ Is connectivity between Jira and Jenkins responding

SDLC Pipeline is Production

- ↘ Monitoring production tells us
 - >>> that our client is satisfied
 - >>> that we are meeting our reliability targets
 - >>> that we are meeting our objectives
 - >>> what's left in the error budget
- ↘ Monitoring the SDLC tells us
 - >>> that we can satisfy the clients need for new features
 - >>> that we are ready to create new releases or deploy
 - >>> how good our burn-down rate is for releasing new features
 - >>> how well our coding is meeting requirements



Client satisfaction

Monitoring Automation

Treat your configuration as code

- >>> Storing in source control management are common practices, with obvious benefits
 - ~ Change history
 - ~ Links from specific changes to tracking system
 - ~ Easier rollbacks and linting checks
 - ~ Enforced code review procedures
 - ~ Intent based preferable to web UI's or CRUD style API's
 - ~ Automated deployment

Encourage consistency

- >>> Centralised approach provides consistency
 - ~ Consistent set of basic metrics means you can
 - Automatically collect metrics across entire organisation provide a consistent set of dashboards
 - Any new component launched has basic monitoring
 - Everyone can access and use monitoring data not just engineers
- >>> Single framework
 - ~ Enables engineers to ramp up faster when they switch teams
 - ~ Easier collaboration during debugging
 - ~ Easily understand another team's dashboard


Monitoring Automation

- ↘ Encourage re-use
- ↘ Links to documentation to help resolve issues on alerts
- ↘ Segregation of duty
 - >>> System monitoring code
 - >>> Application monitoring code
- ↘ System monitoring is common
 - >>> Configurable attributes for common use
 - >>> Infrastructure and platforms responsible for code
- ↘ Application monitoring is specific
 - >>> Enable developer to include monitoring in their SCM
 - >>> Code linked to central monitoring service
 - >>> Developer responsible for application monitoring

Distributed Monitoring


- ↘ Lots of interconnected components
- ↘ Managed by different teams
- ↘ Issues include
 - >>> Time zones (follow the sun)
 - >>> Synchronization
 - >>> Differing versions
 - ~ Operating systems, software components, application
 - >>> Permissions
 - ~ Administrative access may be required to view metrics or logs

Distributed Monitoring – 4 Golden Signals




Latency

- Measure of time to complete action
 - Processing, Response, Round-trip
- Build holistic model of different performance characteristics
- Finding bottlenecks, resources causing delays, success or unsuccessful requests




Traffic

- Capture load and demand for services
- Understand system performance
- Sustained high/low values to determine more resources or route issues
- Where data is moving



Errors


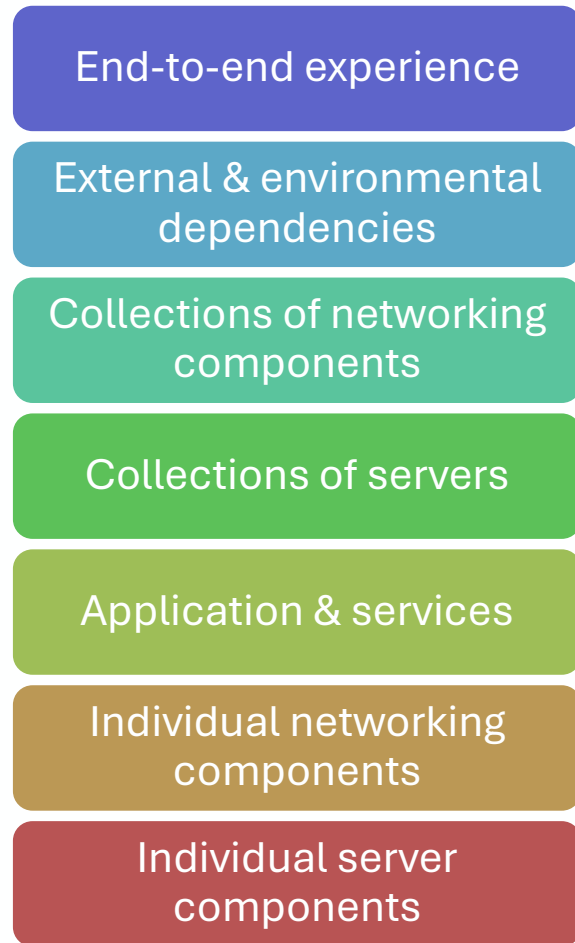
- Component health
- Determine different error types
- Identify root cause



Saturation

- Which resource is being over used – is our load balancing rule effective
- Is our memory clean up policy working
- Saturation could cause traffic delays between different layers

Levels of Abstraction in a System



Increase in
Scope & Level


Tools that deal with distribution

- ↘ AppDynamics
 - >>> Also Cisco now
- ↘ Cisco's Dynatrace
- ↘ Datadog
- ↘ These monitor the interconnectivity, not just the app or O/S

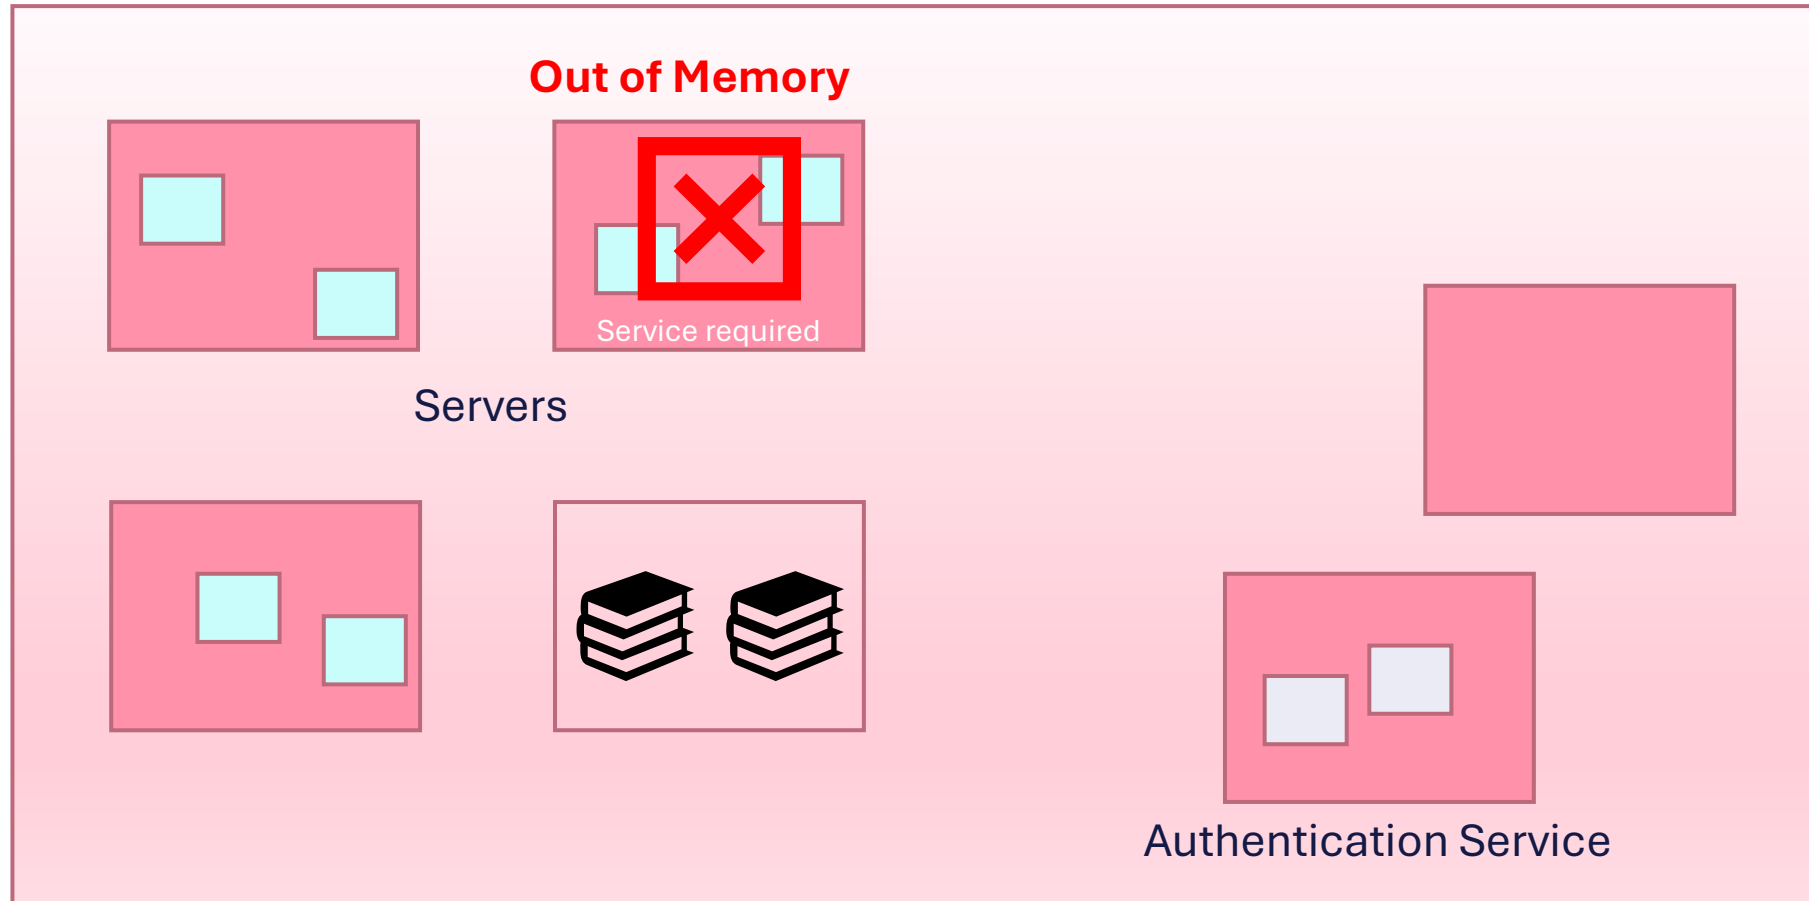
Distributed application tracing

- ↘ AKA Distributed request tracing
 - >>> Ability to track and observe service requests
 - ~ Collection of data as requests go from one service to another
 - ~ Helps to understand flow of requests spanning your operations being traced
- ↘ Especially used for Microservices
 - >>> Pinpoint where failure occur
 - >>> What causes poor performance


Distributed Request Tracing



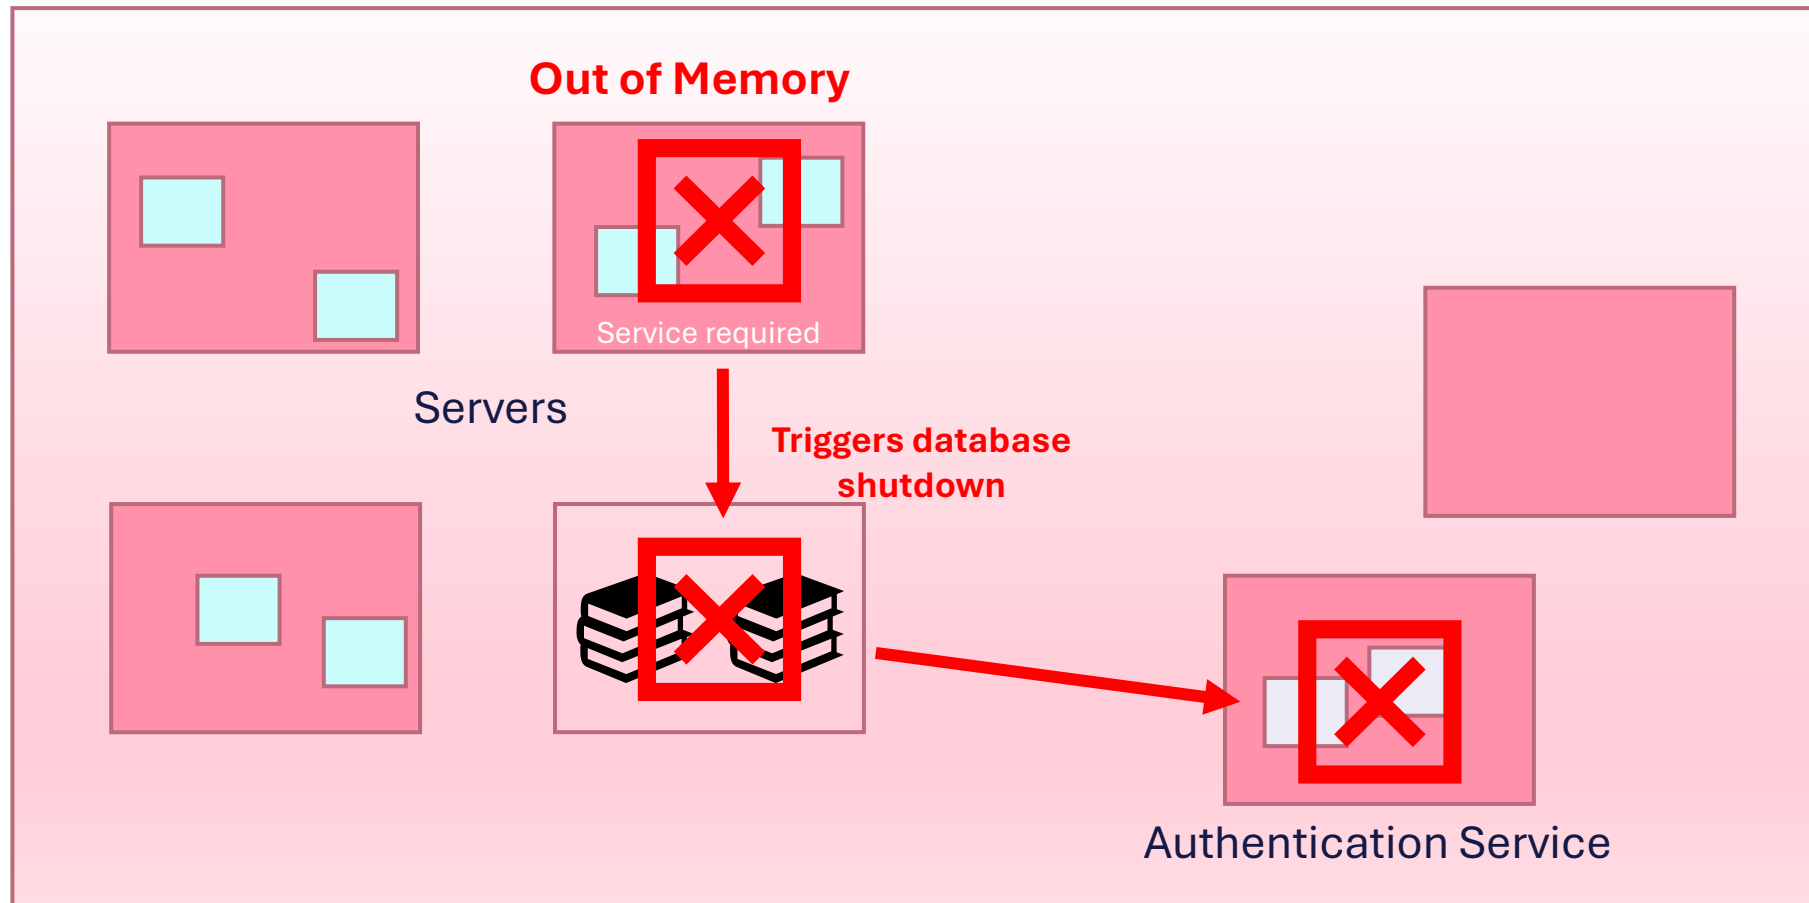
Distributed Request Tracing



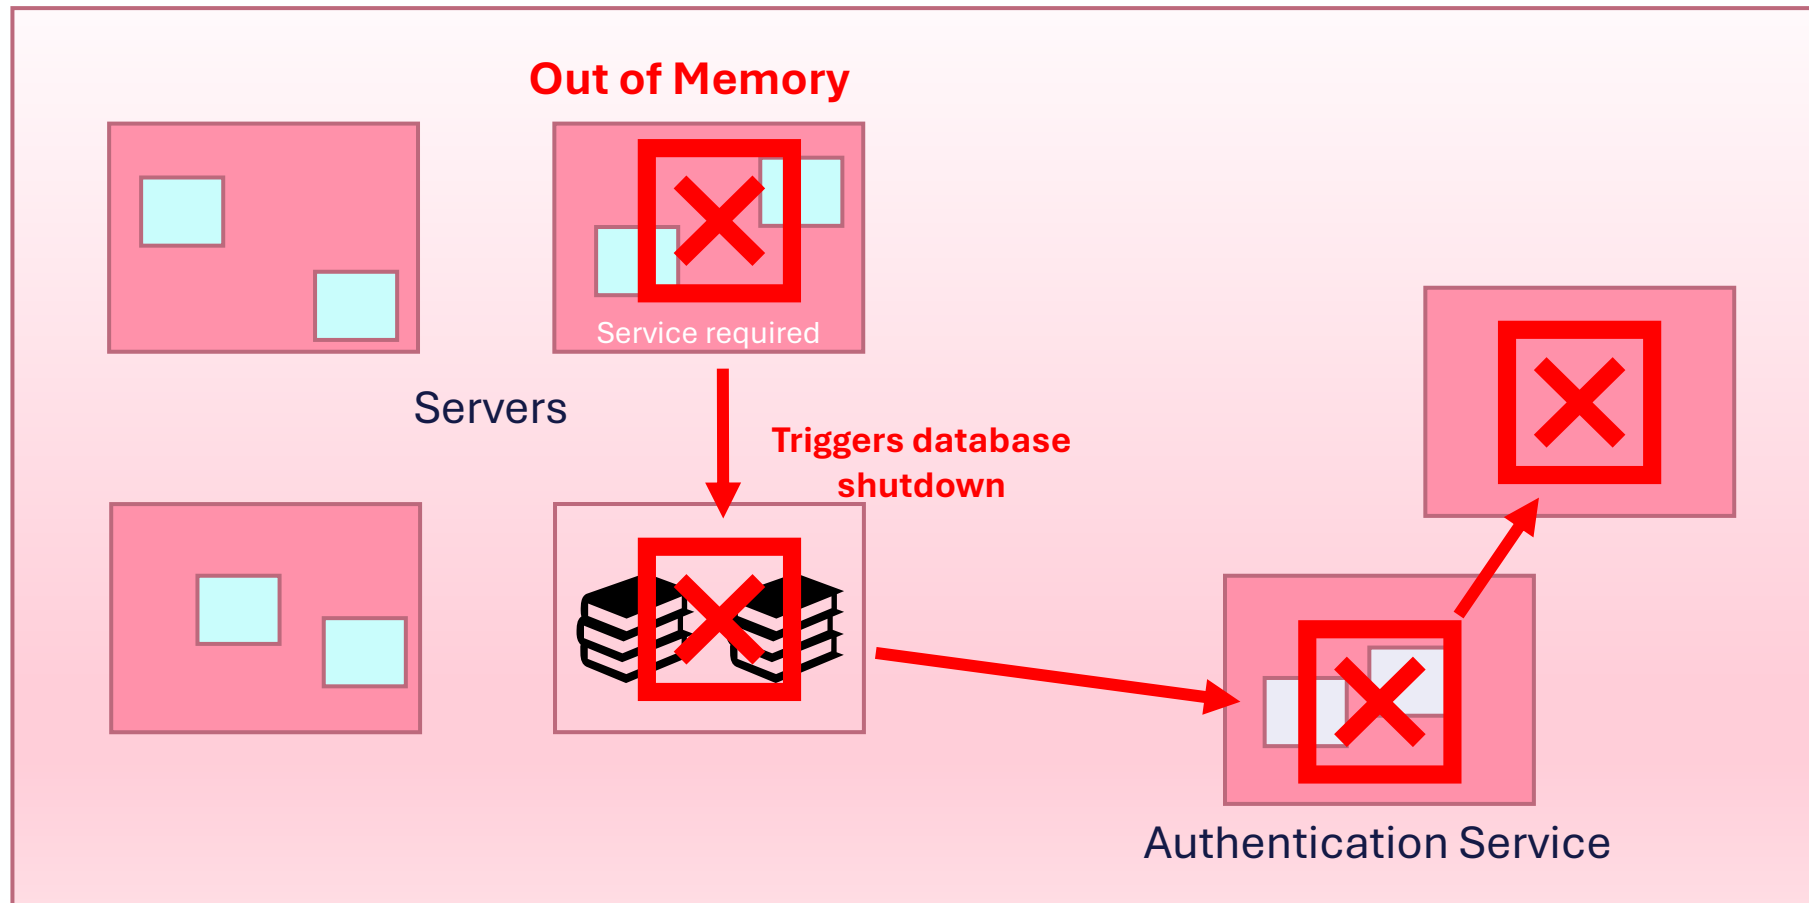
Distributed Request Tracing



Distributed Request Tracing

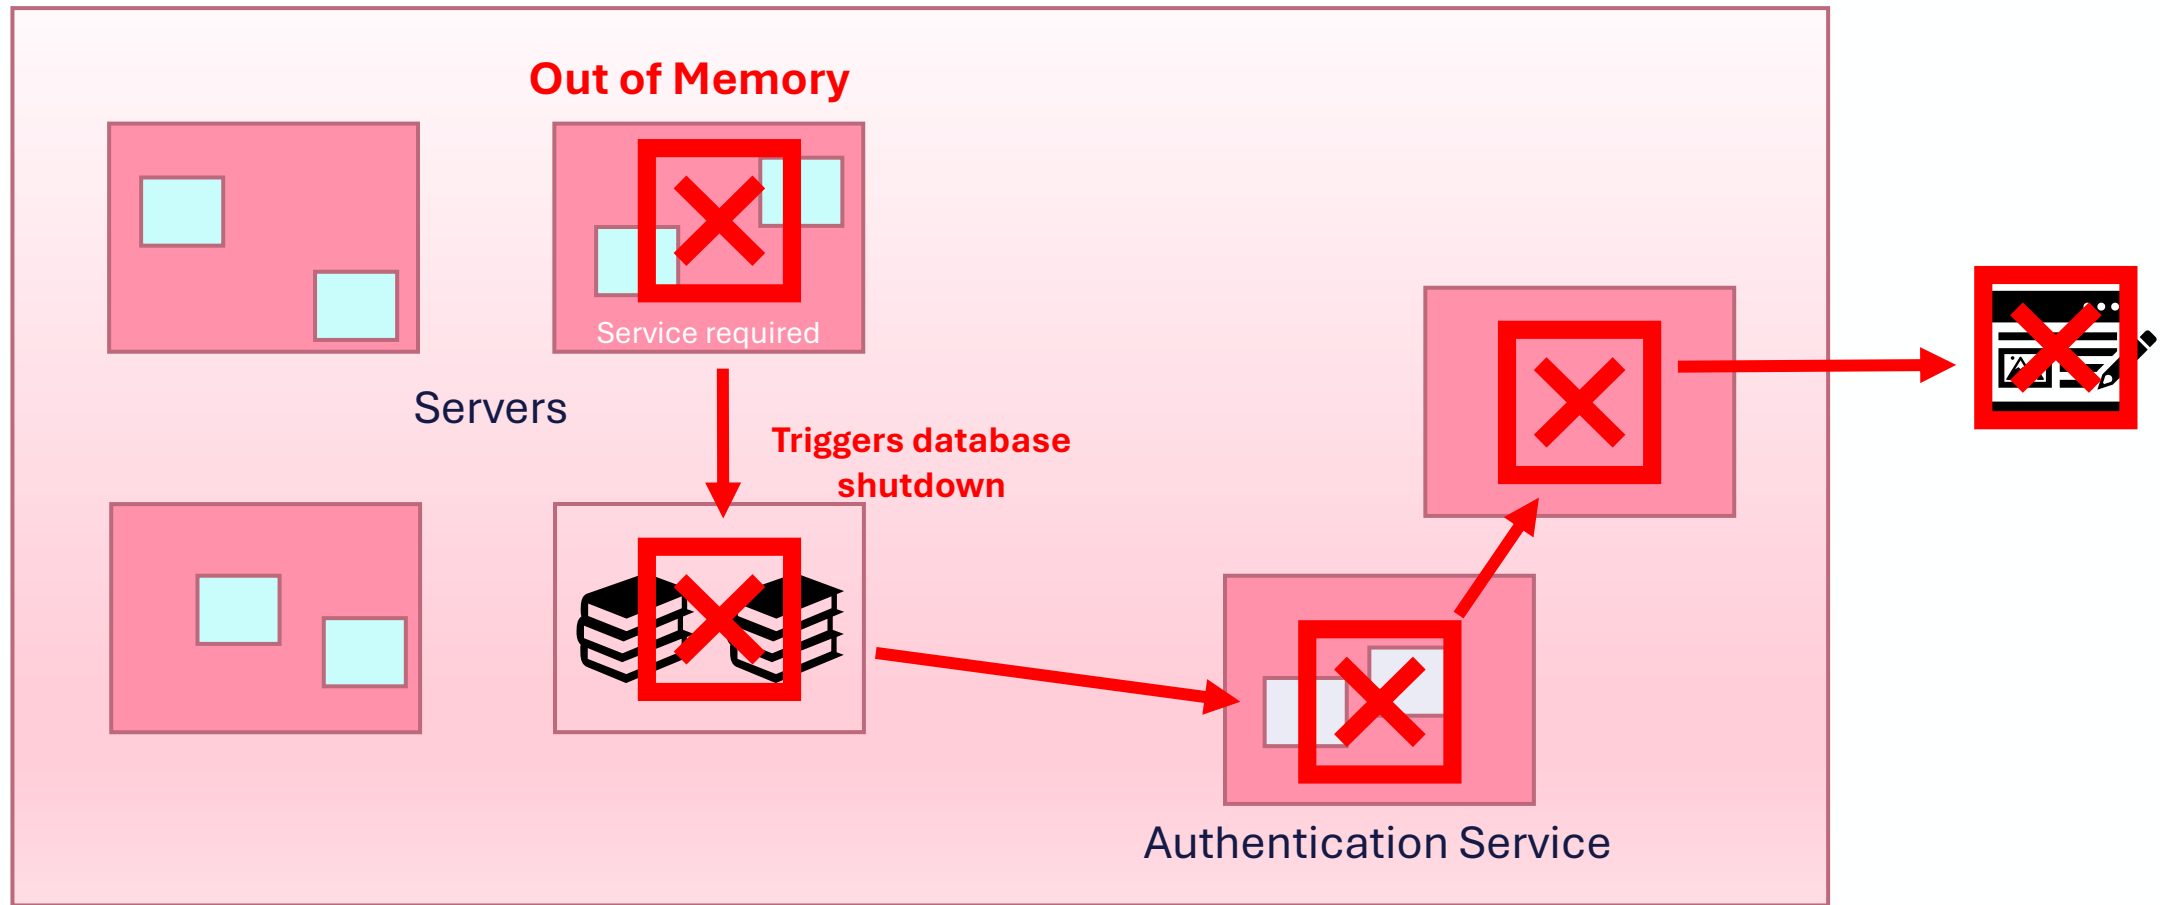


Distributed Request Tracing



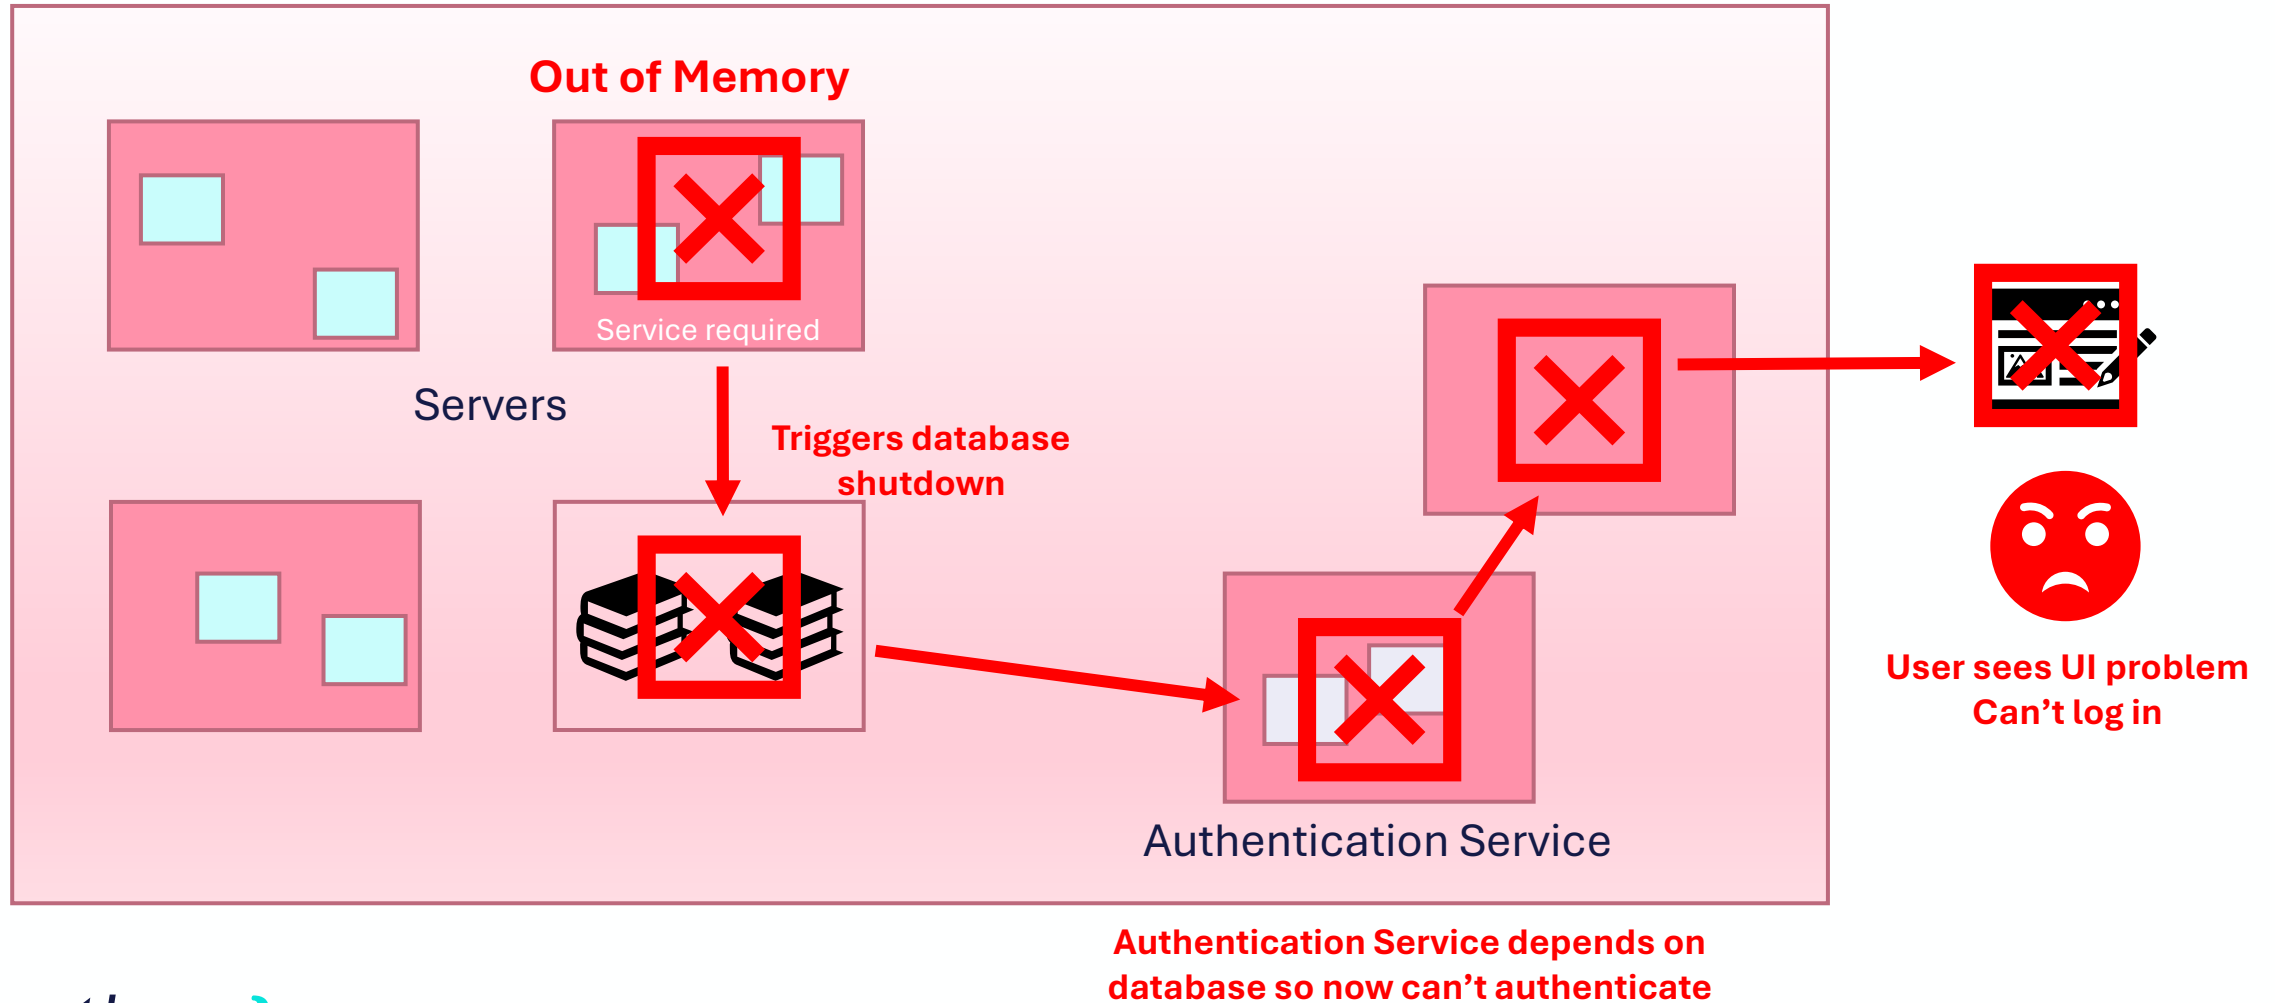
Authentication Service depends on database so now can't authenticate

Distributed Request Tracing



Authentication Service depends on database so now can't authenticate

Distributed Request Tracing



Distributed Tracing

- ↘ Deep understanding of performance of every service
 - >>> Up and down stream
- ↘ Identify and resolve issues to minimize the impact on the customer
- ↘ Measure overall system health
- ↘ Understand the effect of changes on the customer experience
- ↘ Prioritize high-value areas for improvement
- ↘ Continuously improve monitoring to capture new trends
 - >>> Dynamic monitoring systems that detect up/down stream connectivity

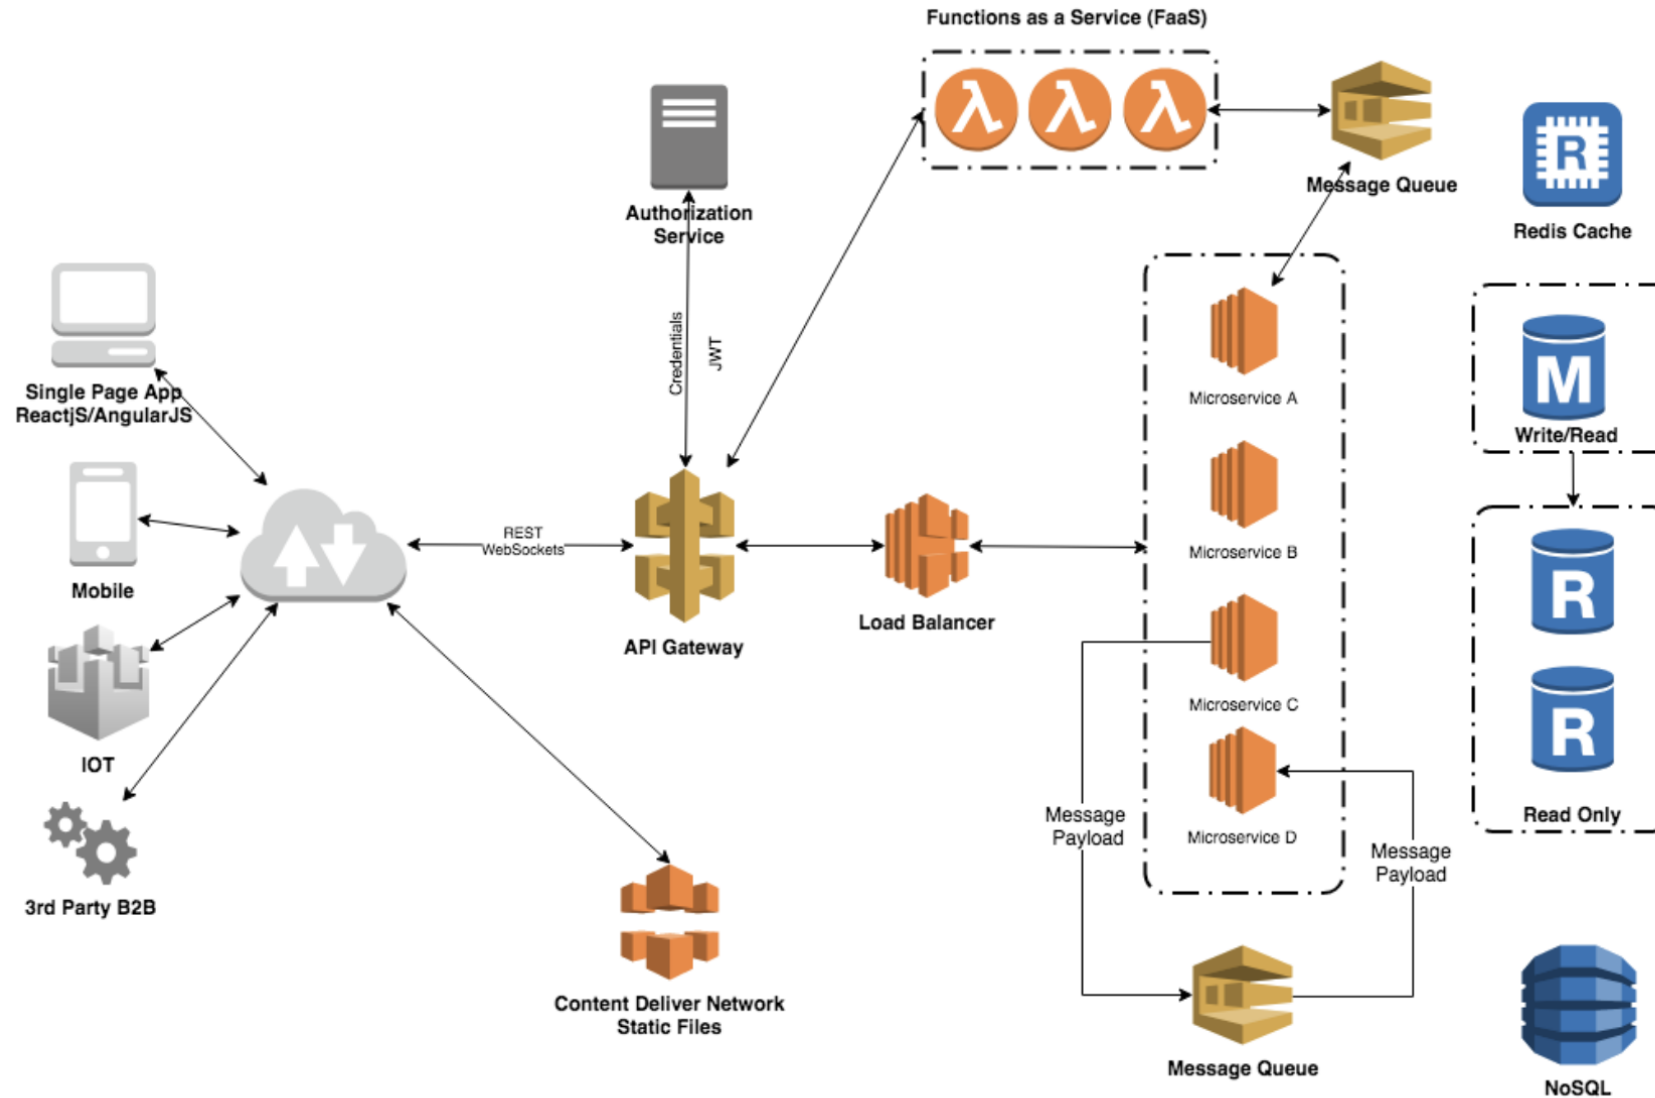
Distributed Tracing

- ↘ Instrumenting environment
- ↘ Trace context
 - >>> Assigns a unique ID to each request
 - >>> Correlates each step in the correct order to allow tracking and monitoring
- ↘ Metrics and metadata
 - >>> Captures data about
 - ~ Spans, errors, duration , Custom attributes
- ↘ Analysis and visualization
 - >>> Provides context needed to derive meaning and assess action

Distributed Tracing

- ↘ Use distributed tracing to get answers to questions such as:
 - >>> What is the health of the services that make up a distributed system?
 - >>> What is the root cause of errors and defects within a distributed system?
 - >>> Where are performance bottlenecks that impact customer experience?
 - >>> Which services have problematic or inefficient code and need prioritizing?

Activity: Distributed Tracing



Monitoring Bottlenecks

- ↘ Avoid overloading the monitoring server
 - >>> Too many systems
 - ~ Loss of important metrics
 - ~ Loss of monitoring data
 - >>> Adjust polling times
 - ~ Spread the times across groups of servers
 - >>> Introduce pyramid system of proxy servers
 - ~ Groups of servers send to a proxy
 - ~ Proxy forwards data for hosts to central server
- ↘ Separate data retrieval from data store
 - >>> Introduce read only servers
 - ~ Dashboards and views use data from read-only replica
 - >>> Use sharded storage services
 - ~ Data spread across multiple server
 - ~ Akin to striping data across disks, but across servers instead

- ↘ Separate the service components
 - >>> Dashboard
 - >>> Storage
 - >>> Alerting
 - >>> Collecting



Rules for Effective Monitoring Management

- ↘ As simple as possible
- ↘ Avoid piling up the requirements
 - >>> Leads to complex monitoring systems
 - >>> Complexity introduces
 - ~ Differing latency thresholds
 - ~ Different percentiles on different kinds of metrics
 - ~ Specific dashboard components for each type of cause
 - >>> Complexity increases with time
 - ~ Monitoring system becomes fragile, difficult to change, increase in toil
- ↘ Design with simplicity in mind
 - >>> Rules to catch real incidents – simple, predictable, reliable
 - >>> Rarely used data collection and aggregation should be removed
 - >>> Rarely used features on dashboards should be removed
- ↘ SRE toil reduction methods should be applied

Summary

- ↘ SREs to be familiar with a service's monitoring system and features
- ↘ SREs require monitoring to define users experience of service
- ↘ Need to know
 - >>> Where to look
 - >>> How to identify abnormal behaviour,
 - >>> How to find the information they need during an emergency
- ↘ Combine some source of metrics and logging in your monitoring strategy
 - >>> Exact mix is highly context-dependent
 - >>> Collect metrics that serve a particular purpose
 - ~ Better capacity planning
 - ~ Assist in debugging
 - ~ Directly notify you of problems



Q&A

Reference

↘ Mushero, S. (2019, August 2). Push vs. Pull Monitoring Configs. Retrieved from <https://steve-mushero.medium.com/push-vs-pull-configs-for-monitoring-c541eaf9e927>