



Observability

Site Reliability Engineering



Overview

In this module, we will provide an overview of the concept of observability.

Learning Objectives

- ↘ Describe the parts of observability
- ↘ Discuss the concept of application performance monitoring
 - >>> Business transactions
 - >>> Component monitoring
- ↘ Explain the relationship between service level indicators (SLIs), service level objects (SLOs), and service level agreements (SLAs)

Observability

- ↘ More than monitoring
- ↘ Three parts
 - >>> Logs
 - >>> Metrics
 - >>> Traces
- ↘ Using data from a complex system to infer its internal state
 - >>> Capacity & Performance
 - >>> Satisfaction
 - >>> Expectations

Three Parts of Observability

1. Logs

>>> Record of past events

2. Metrics

>>> Current data about the system components

3. Traces

>>> Capture activity for a business transaction

>>> Shows interactivity in complex systems

Log Files

- ↘ Automatically created
 - >>> Application
 - >>> Operating system
- ↘ Hold information
 - >>> User behaviour
 - >>> Events
- ↘ Used for root cause analysis
 - >>> Understand why a metric changed
 - >>> Identify where an event began




Metrics

- ↘ Quantified measurement
 - >>> Used to understand the status of a specific process
 - >>> Often compared to a defined baseline
 - ~ Analyse the system's or process's status
- ↘ Trends in metric changes
 - >>> Indication of underlying issue



Trace

- ↘ Complete record of business request
 - >>> Illustrates a complete transaction
 - >>> Captures all the components and services involved
- ↘ Contains hundreds of data points that can:
 - >>> Indicate errors
 - >>> Diagnose security threats
 - >>> Detect and isolate component or network issues




Observability Data

- ↘ Combining three parts to answer questions
 - >>> Why performance is degrading
 - >>> What dependency behaviors have changed
 - >>> Why this application is failing
 - >>> Where to look for a fix



Expectations

- ↘ Focus on user expectations
 - >>> Delivering more than expected doesn't gain much
 - >>> Failing to deliver will lose a lot
- ↘ Observability should provide at a glance:
 - >>> Are we out of SLO?
 - >>> Should SLO be adjusted?
 - >>> Where is the problem?



Benefits of Observability

- ↘ Overview of complex system
- ↘ Faster feature release
- ↘ Observe impact of updates
 - >>> Confirmation improvements
 - >>> Identify issues

Application Performance Monitoring

- ↘ Purpose of application monitoring
- ↘ Components of APM

Purpose of Application Monitoring

- Essential to maintain uninterrupted business processes
- Monitor the overall application
 - >>> Continued availability
 - >>> Appropriate performance
- APM solution provides:
 - >>> Connection of app performance to business outcomes
 - >>> Isolate and fix errors
 - ~ Before they affect the end user
 - >>> Reduce the mean time to repair (MTTR)



Five Most important Elements

1. Runtime application architecture discovery
2. End-user experience monitoring
3. User-defined transaction profiling
4. Component deep-dive monitoring
5. Analytics

Runtime Application Architecture

- ↘ Analyses the hardware and software components used
 - >>> Direct execution
 - >>> To communicate with
- ↘ Anticipate potential problems
 - >>> Pattern recognition
 - >>> Performance problems



Real User Monitoring

- ↘ Real user monitoring enables an organisation to efficiently respond to faults and understand their impact
 - >>> Also referred to as *end-user experience monitoring*
- ↘ Gather user-based performance data
 - >>> How well the application is performing (from user point of view)
 - >>> Gauge potential performance problems
- ↘ Performance through entire infrastructure
- ↘ Details on the analysed client
 - >>> Location
 - >>> Operating system
 - >>> Browser

Two Ways to Track End User Experience

Synthetic Monitoring

- ↘ Uses probes and bots
- ↘ Simulates an end user to determine problems before the app is opened
- ↘ Used to monitor service-level agreements

Agentless monitoring

- ↘ Uses data probes
- ↘ Analyse network traffic that travels
 - >>> Load balancers
 - >>> Switches

Business Transactions

Focus on specific user interactions

- ↘ Recreating them to test and understand the conditions that lead to a performance problem
- ↘ AKA - User-defined transaction profiling

Help organizations

- ↘ Trace business transaction movement across various components
- ↘ Reveal when and where events are occurring
- ↘ Optimize performance by identifying bottlenecks

Component Monitoring

- ↘ Provides a deeper understanding of the specific elements and pathways
- ↘ Also referred to as an application component deep dive
 - >>> Tracking all components of the IT infrastructure
- ↘ Extensive, in-depth monitoring
 - >>> All resources and events
 - ~ Analysis of all servers
 - ~ Operating systems
 - ~ Middleware
 - ~ Application components
 - ~ Network components

Analytics and Reporting

- ↘ Essential to ensuring the organization receives a good return on investment
- ↘ Translating data gathered into information that can be used
 - >>> Define a performance baseline using historical and current data
 - ~ Set expectations for normal app performance
 - >>> Identify areas of improvement
 - ~ Comparing infrastructure changes to performance changes
 - >>> Identify, locate and resolve performance issues
 - ~ Using historical and baseline data
 - >>> Predict and alleviate potential future issues
 - ~ Before the customer notices

Critical APM Metrics

↘ Web performance monitoring

- >>> Average response time for end user interactions
- >>> Identify if speed is affecting app performance

↘ System metrics impacting app performance

- >>> CPU usage
- >>> Disk read/write speeds
- >>> Memory demands

↘ Application availability and uptime

- >>> App is online and available to users
- >>> Frequently used to determine compliance with SLA

↘ Request rates

- >>> Amount of traffic received by the application
- >>> Identify significant increases, decreases
- >>> Coinciding users

↘ Customer satisfaction

- >>> Compare how customers feel about the app against defined baseline

↘ Error rates

- >>> Capture app degrading or fails at the software level

↘ Number of instances

- >>> How many server or app instances are running
- >>> Important for cloud application

Observing Toil

- ↘ How do we measure toil success?
- ↘ Where do we get the metrics?



Observing Toil

- ↘ Toil backlog
 - >>> Should be reducing
 - >>> Should not see the same toil recurring in the backlog
 - >>> Amount of toil being reported by specific people reducing
 - >>> Error budget not decreasing as fast = more reliable system
- ↘ Reduced fatigue in the team
- ↘ Shorter MTTR
- ↘ Toil metrics come from
 - >>> Ticketing and job systems such as Jira and ServiceNow
 - >>> Other systems that gather information about time at work



SLAs vs SLOs vs SLIs

Service level agreements

Formal financial or contractual agreements
Usually with penalties

Service level objectives

Reflect customer expectations
Based on customer satisfaction
Majority of customers happy
Always outliers

Service level indicators

Show progress towards SLO
Indicate health of business

Summary Q & A

- ↘ Observability
 - >>> Logs
 - >>> Metrics
 - >>> Traces
- ↘ Application Performance Monitoring
 - >>> Business Transactions
 - >>> Component Monitoring
- ↘ Service level indicators indicate service level objects which meet service level agreements