



SRE in Action

Site Reliability Engineering



Overview

- ↘ DevOps vs SRE
- ↘ SRE Activities
- ↘ Advantages of SRE
- ↘ Error Budget & Calculations
- ↘ 7 Principles in detail

DevOps vs. SRE

DevOps

- ↘ The intersection of two key disciplines
 - >>> Software Development (Dev) and IT Operations (Ops).

SRE

- ↘ The application of software engineering to operational problems
 - >>> Teach application developers how to build reliable services

Typical Activities

- ↘ Develop and manage scalable, secure and stable systems
- ↘ Conduct Incident analysis
- ↘ Analyse performance and create improvement plans
- ↘ Monitor system efficiency
- ↘ Manage risks
- ↘ Automate manual tasks within the SDLC
- ↘ Ease workload through automated tools, logs and testing environments
- ↘ Implement new features
- ↘ Select infrastructure tools
- ↘ Adapt environments to increasing or decreasing numbers of users

SLIs versus SLOs

SLIs (Service Level Indicators)

- ↘ Real metrics obtained from monitoring systems and other systems that can provide metrics
- ↘ These metrics are based on values that will impact end-user experience

SLOs (Service Level Objectives)

- ↘ Agreed goals by the project team (Devs, Ops, Management)
- ↘ SLIs determine the basis for the SLO goals
- ↘ A tangible set of values to demonstrate to the client our commitment

The Error Budget

- ↘ A tool used to balance service reliability with pace of innovation
- ↘ A control mechanism for diverting attention to stability as needed

Google

Changes are a major source of instability, representing roughly 70% of our outages, and development work for features competes with development work for stability

Error Budget Calculations

↘ Embracing risk and managing risk requires measurement

↘ Time based risk

$$availability = \frac{uptime}{(uptime + downtime)}$$

↘ Aggregate availability

$$availability = \frac{successful\ requests}{total\ requests}$$

An Error Budget

1 minus the SLO of the service

100% reliability is never the right target!

- ↘ A 99.9% SLO service has a 0.1% error budget
- ↘ Example:
 - >>> If our service receives 1,000,000 requests in four weeks,
 - >>> A 99.9% availability SLO gives us a budget of 1,000 errors over that period
- ↘ Managing service reliability is largely about managing risk
- ↘ An error budget aligns incentives and emphasizes joint ownership between SRE and product development
- ↘ Error budgets help decide rates of releases and effectively defuse talks about outages with stakeholders
- ↘ Allows multiple teams to reach the same conclusion about production risk without rancour

Activity: Calculate an Error Budget

- ↘ The app devs state that to keep a customer happy, they need the DBAs to ensure that the databases will
 - >>> Serve 10,000,000 requests per working day (Mon – Fri)
 - >>> Time to service requests 1ms
 - >>> Over a 4-week period, we can cope with a maximum delay of 5 seconds
- ↘ What is our Error Budget as a percentage?
- ↘ What is our reliability as a percentage?

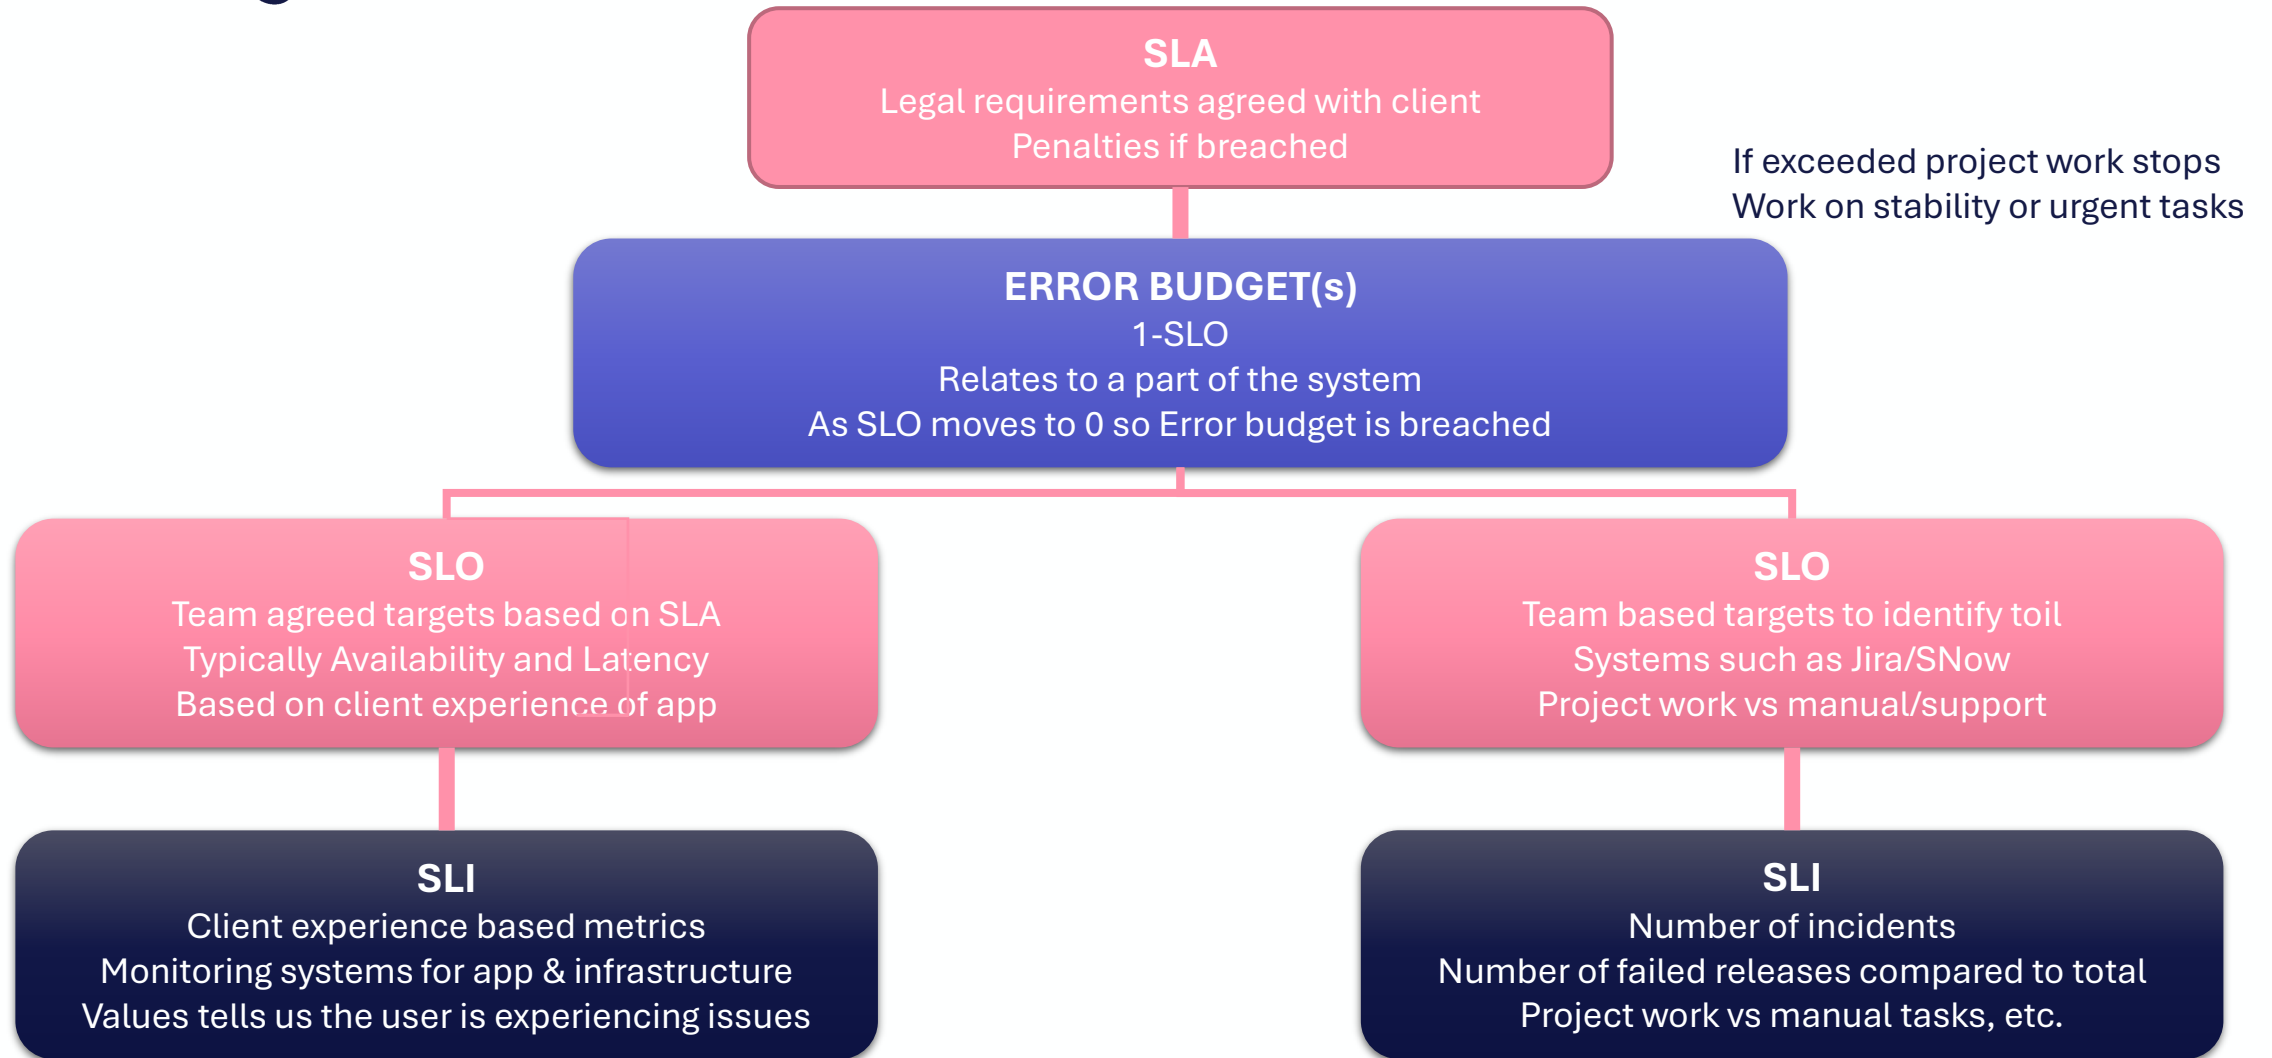
Activity: Calculate an Error Budget (Answer)

- ↘ The app devs state that to keep a customer happy, they need the DBAs to ensure that the databases will
 - >>> Serve 10,000,000 requests per working day (Mon - Fri)
 - ~ 200,000,000 requests in a 4-week period
 - >>> Time to service requests 1ms
 - ~ 200,000,000 ms in a 4-week period
 - >>> Over a 4-week period, we can cope with a maximum delay of 5 seconds
 - ~ $5 \times 1000 = 5000$
- ↘ $(10,000,000 \times 20) / (5 \times 1000) = 40,000$ ms in a 4-week period
- ↘ Error budget 0.02% = 99.98% reliability

SRE 7 Principles in Detail

1. Embracing risk
 - >>> Manage, be open and learn
 - >>> Every release comes with risk
 - >>> Be prepared
2. Service level objectives
 - >>> Decide as a team = SRE, DevOps, Product Owners
3. Eliminating toil
 - >>> If it doesn't need human decisions, automate it
 - >>> If it costs more to automate than to perform the task, leave it
4. Monitoring distributed systems
 - >>> Reduce fatigue
 - >>> What broke, when and why
 - >>> Fix root causes to prevent repeats
5. Automation
 - >>> The move to self-healing systems
 - >>> Anything that doesn't need human decision or interaction
 - >>> When it doesn't exceed the error budget or the total time to fix manually
6. Release engineering
 - >>> Consistent deployment
 - >>> Solid understanding of SCM, Testing, CI, Post deployment monitoring
7. Simplicity
 - >>> "At the end of the day, our job is to keep agility and stability in balance"
 - >>> Smaller releases, easier measurements

SRE Diagram





Summary Q & A

Resources

- ✚ Alvidrez, M. (2016). Chapter 3: Embracing Risk. In Site Reliability Engineering: How Google Runs Production Systems, Eds. Beyer, B., Jones, C., Petoff, J, & Murphy, Niall. From <https://sre.google/sre-book/embracing-risk>
- ✚ Beyer, B. et al. (n.d.) Google SRE Workbook. From <https://sre.google/workbook/table-of-contents/>
- ✚ Grams, C. (15 Oct 2019). How Much Time Do Developers Spend Actually Writing Code? From <https://thenewstack.io/how-much-time-do-developers-spend-actually-writing-code/>
- ✚ Vizard, M. (16 Feb 2021). Survey: Fixing Bugs Stealing Time from Development. From <https://devops.com/survey-fixing-bugs-stealing-time-from-development/>